

Introduction to GPU programming

Outline

Day 1

(09:00 – 13:00)

- ❑ Introduction
- ❑ Coding a CUDA kernel

- ❑ Example
 - ❑ Matrix Multiplication + Exercise

- ❑ Real world implementations
 - ❑ Genomes simulation
 - ❑ Random numbers generation

- ❑ Homework

Day 2

(09:00 – 13:00)

- ❑ Recap and homework solutions

- ❑ Common pitfalls

- ❑ Memory management
 - ❑ Image processing example
 - ❑ Sequence alignment

- ❑ Exercise
 - ❑ Area of the Mandelbrot set

- ❑ Monitoring and asynchronicity
 - ❑ Smoothed Particle Hydrodynamics

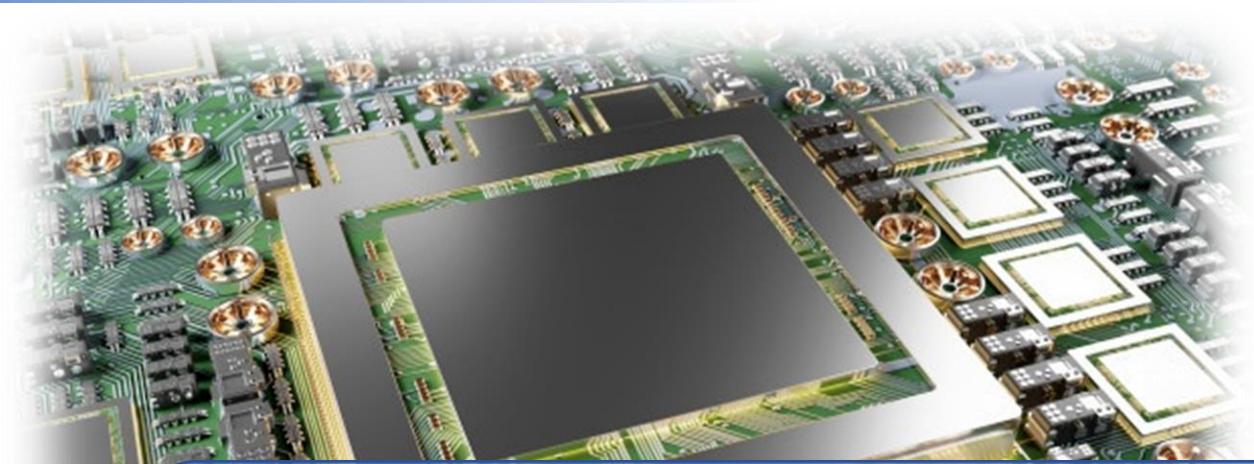
- ❑ Final remarks

Outline

Day 1

(09:00 – 13:00)

- ❑ Introduction
- ❑ Coding a CUDA kernel
- ❑ Example
 - ❑ Matrix Multiplication + Exercise
- ❑ Real world implementations
 - ❑ Genomes simulation
 - ❑ Random numbers generation
- ❑ Homework



CPU

Central Processing Unit

GPUs are designed to execute the same operation in parallel on many independent data elements, while CPUs are designed to execute a single stream of instructions as quickly as possible.

GPU

Graphics Processing Unit

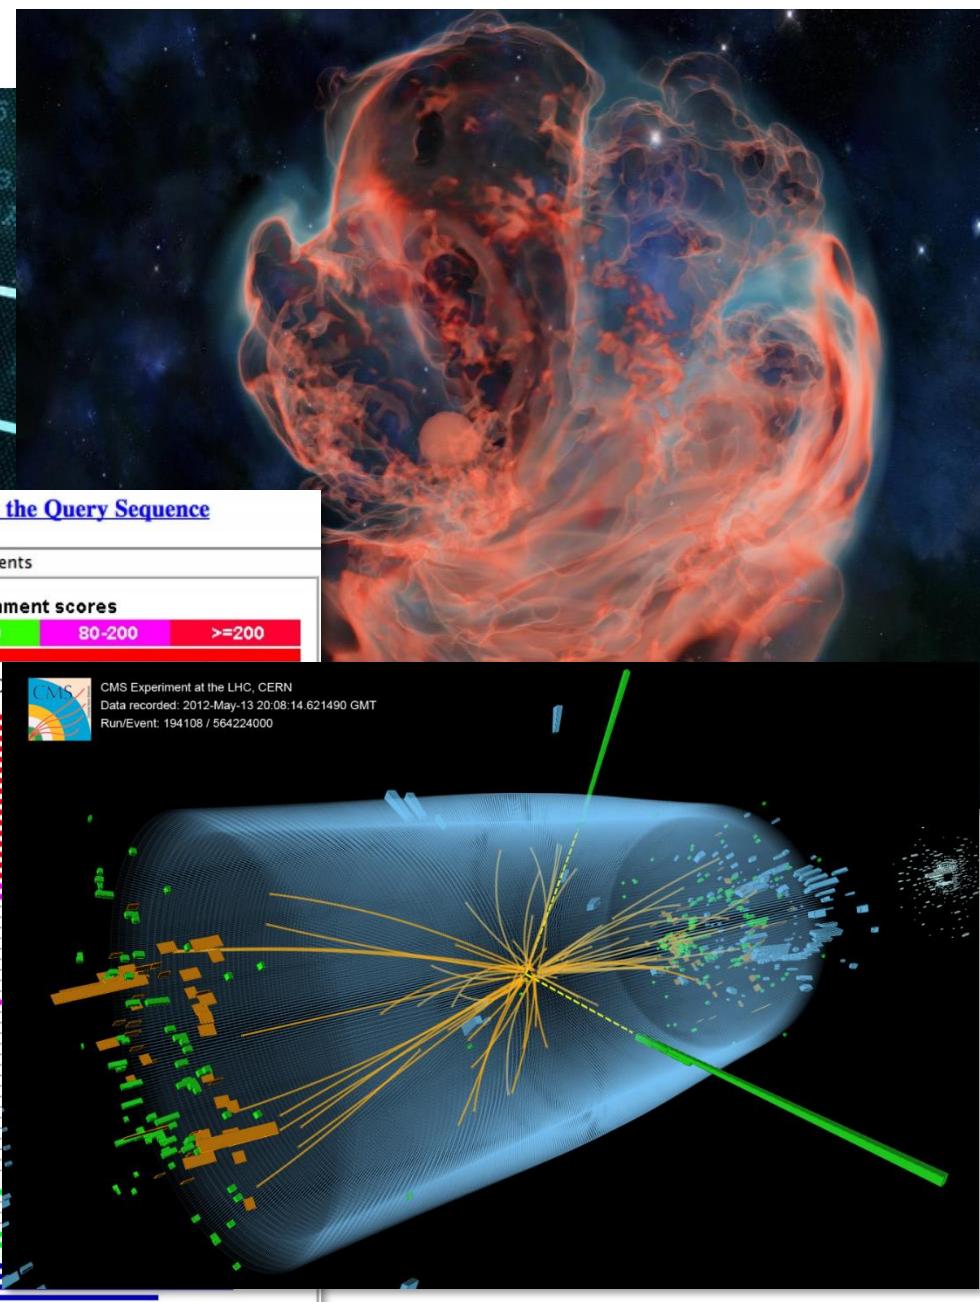
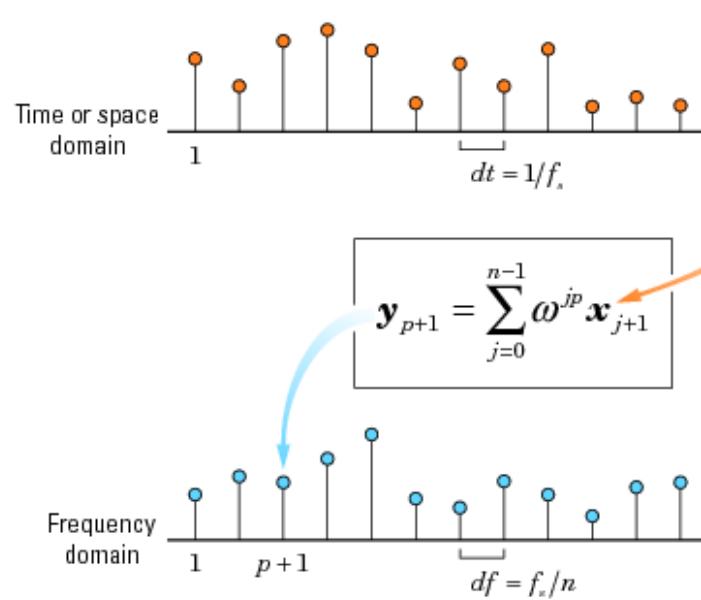




GPU

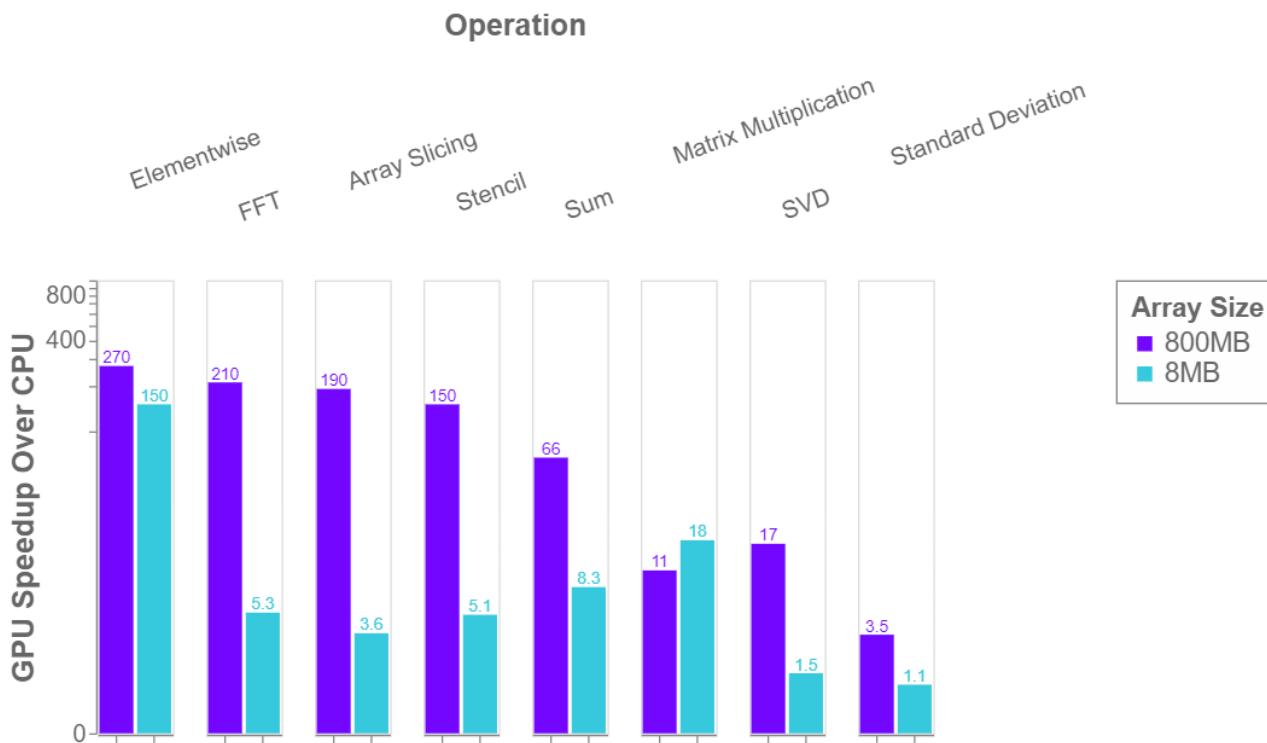
Graphics Processing Unit





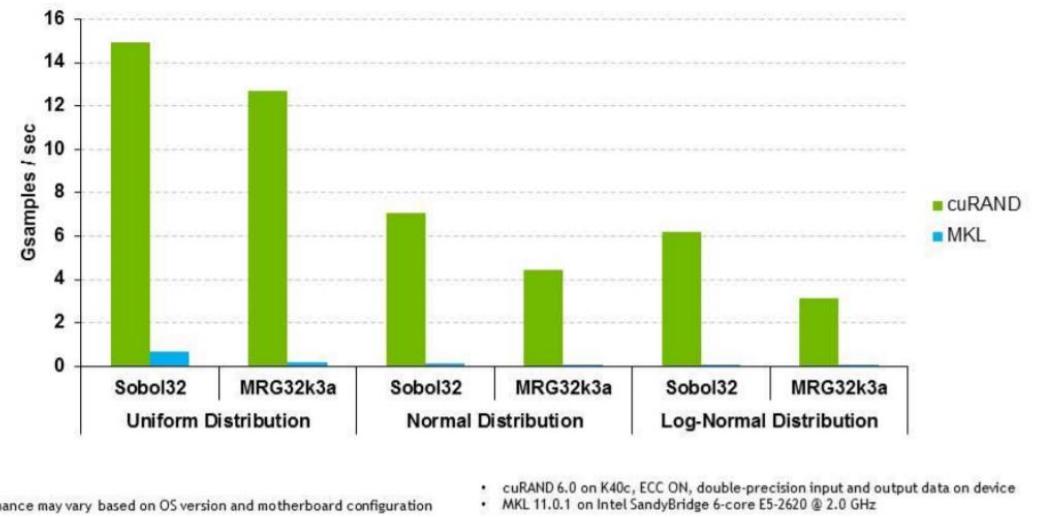
GPGPU
General Purpose
Graphics Processing Unit

Introduction



Single GPU using CuPy library
(GPU optimized library)
compared to numpy

cuRAND: Up to 75x Faster vs. Intel MKL



Comparison between cuRAND
library (GPU optimize) and
MKL library (CPU optimize)

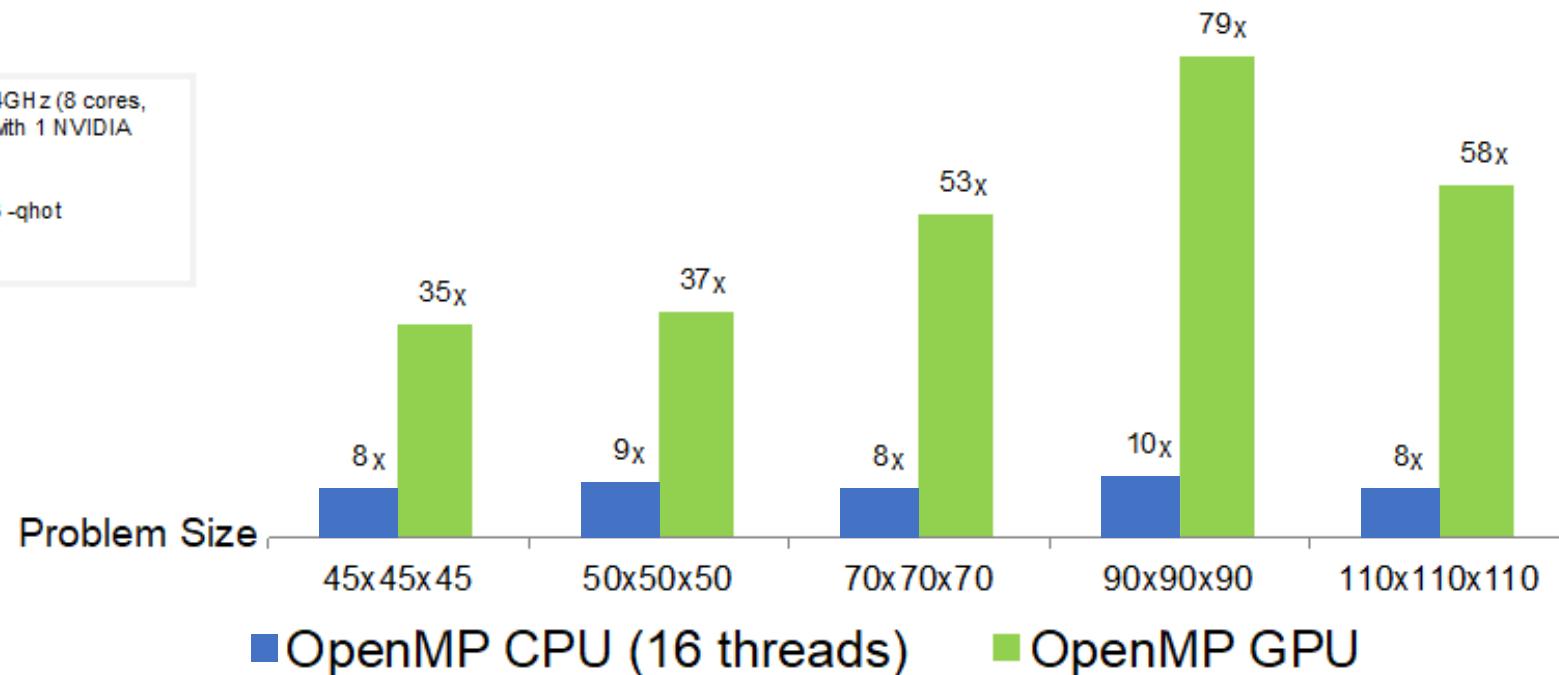
Source: <https://blog.dask.org/2019/06/27/single-gpu-cupy-benchmarks>

http://www.eurorisksystems.com/documents/speed_up_of_numeric_calculations_using_GPU.pdf

Test Specs

2 Power8 sockets @ 4GHz (8 cores, with 8 threads each) with 1 NVIDIA Pascal P100 GPU.

Compiler Options: -O3 -qhot
-qsmp=omp -qoffload*
* Where applicable



IBM Systems

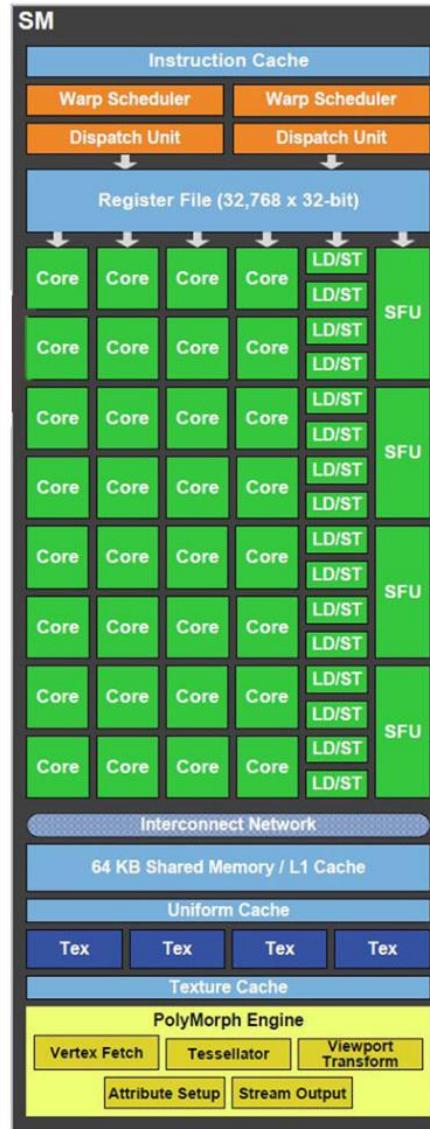
LULESH benchmark changing one single OMP directive.

Source: <https://www.openmp.org/updates/openmp-accelerator-support-gpus/>

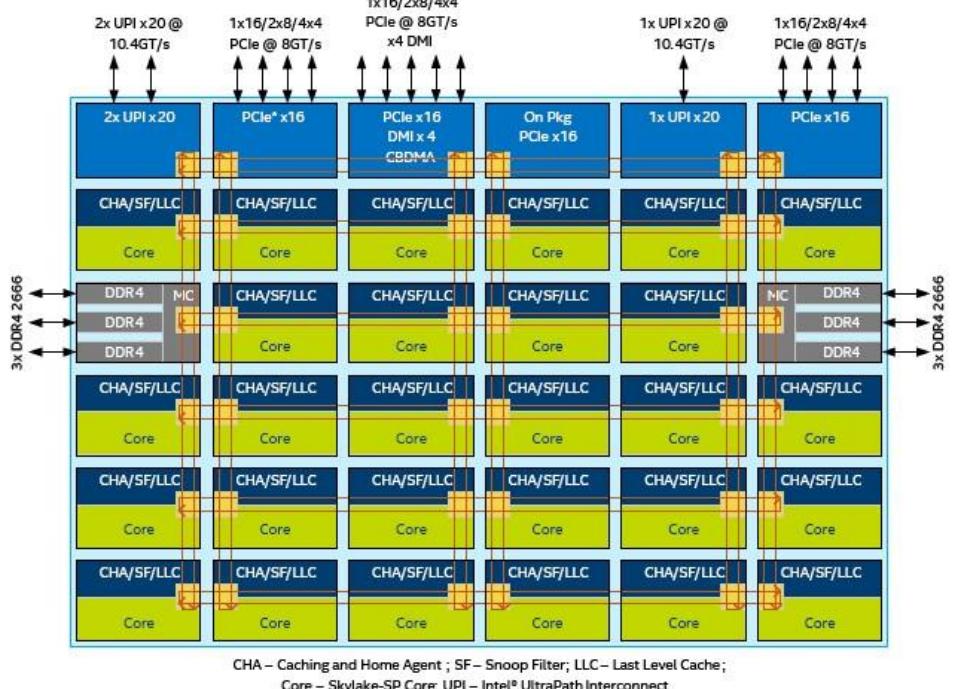
Introduction

Streaming Multiprocessor

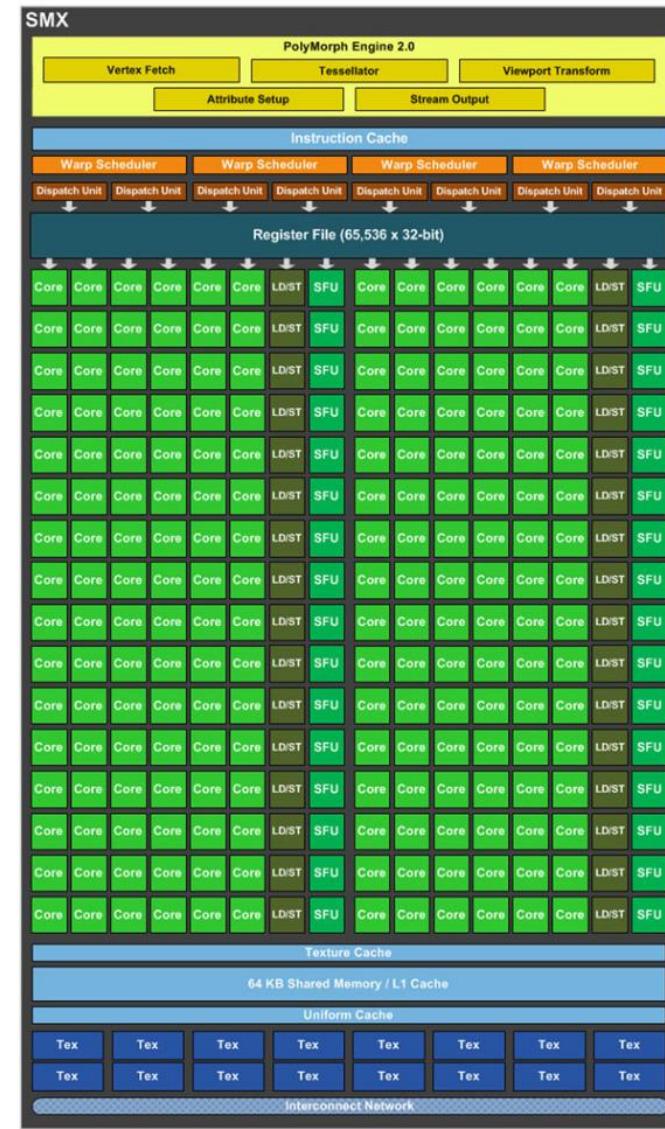
Tesla SM unit (2007)



Intel Skylake (2017)



Fermi SMX unit (2011)



CPU	GPU
Large memory	Relatively small memory
Each core has its own independent control logic <i>Allows independent execution</i>	Groups of cores share control logic <i>Saves space & power</i>
Coherent caches between cores	Shared cache & sync only within groups
Tuned for serial execution of independent work <i>MIMD</i>	Tuned for parallel execution of the same work <i>SIMD</i>
Multiple independent threads	Threads work in lockstep (warp) within groups
It has branch prediction	It serializes codes with branches
Memory latency hidden by cache & prefetching <i>Requires regular data access patterns</i>	Memory latency hidden by scheduling stalled threads <i>Requires 1000s of concurrent threads</i>
Hyperthreading & Vectorization	None
Out of order execution	None

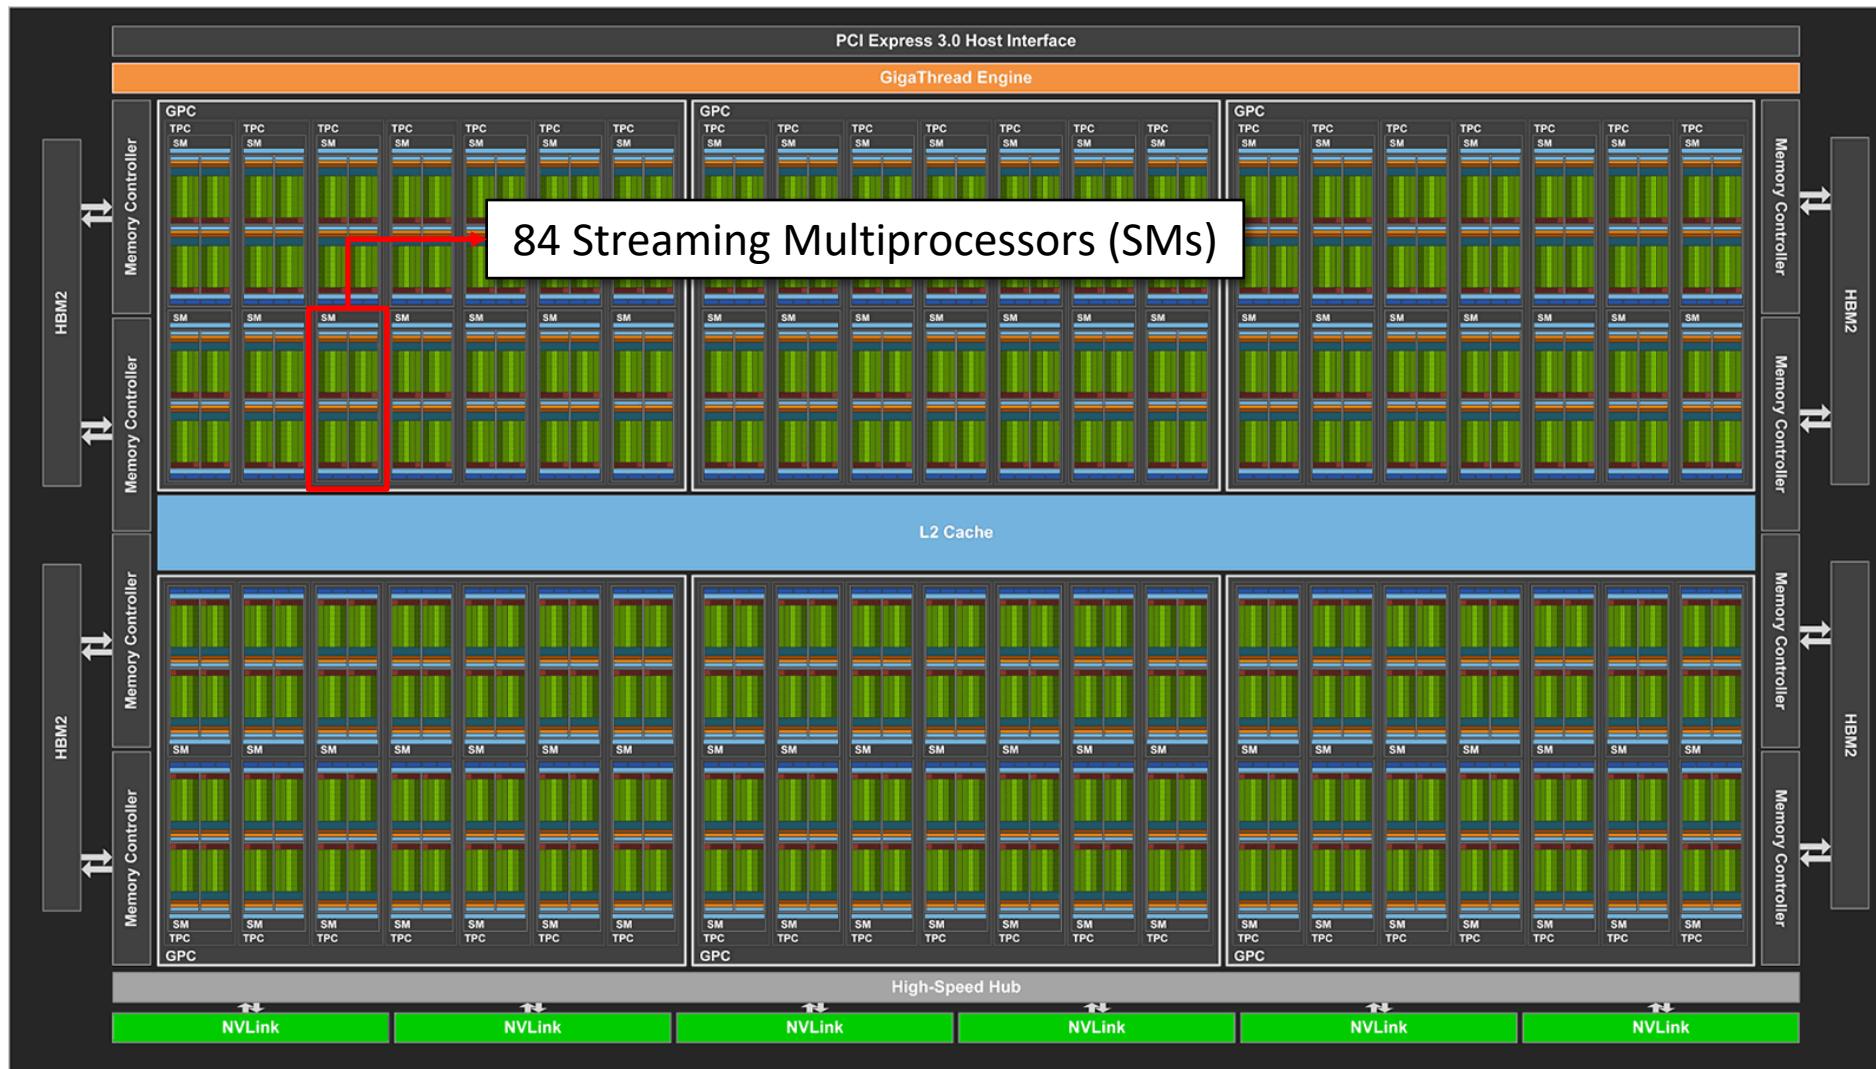
Introduction

Volta (2017)



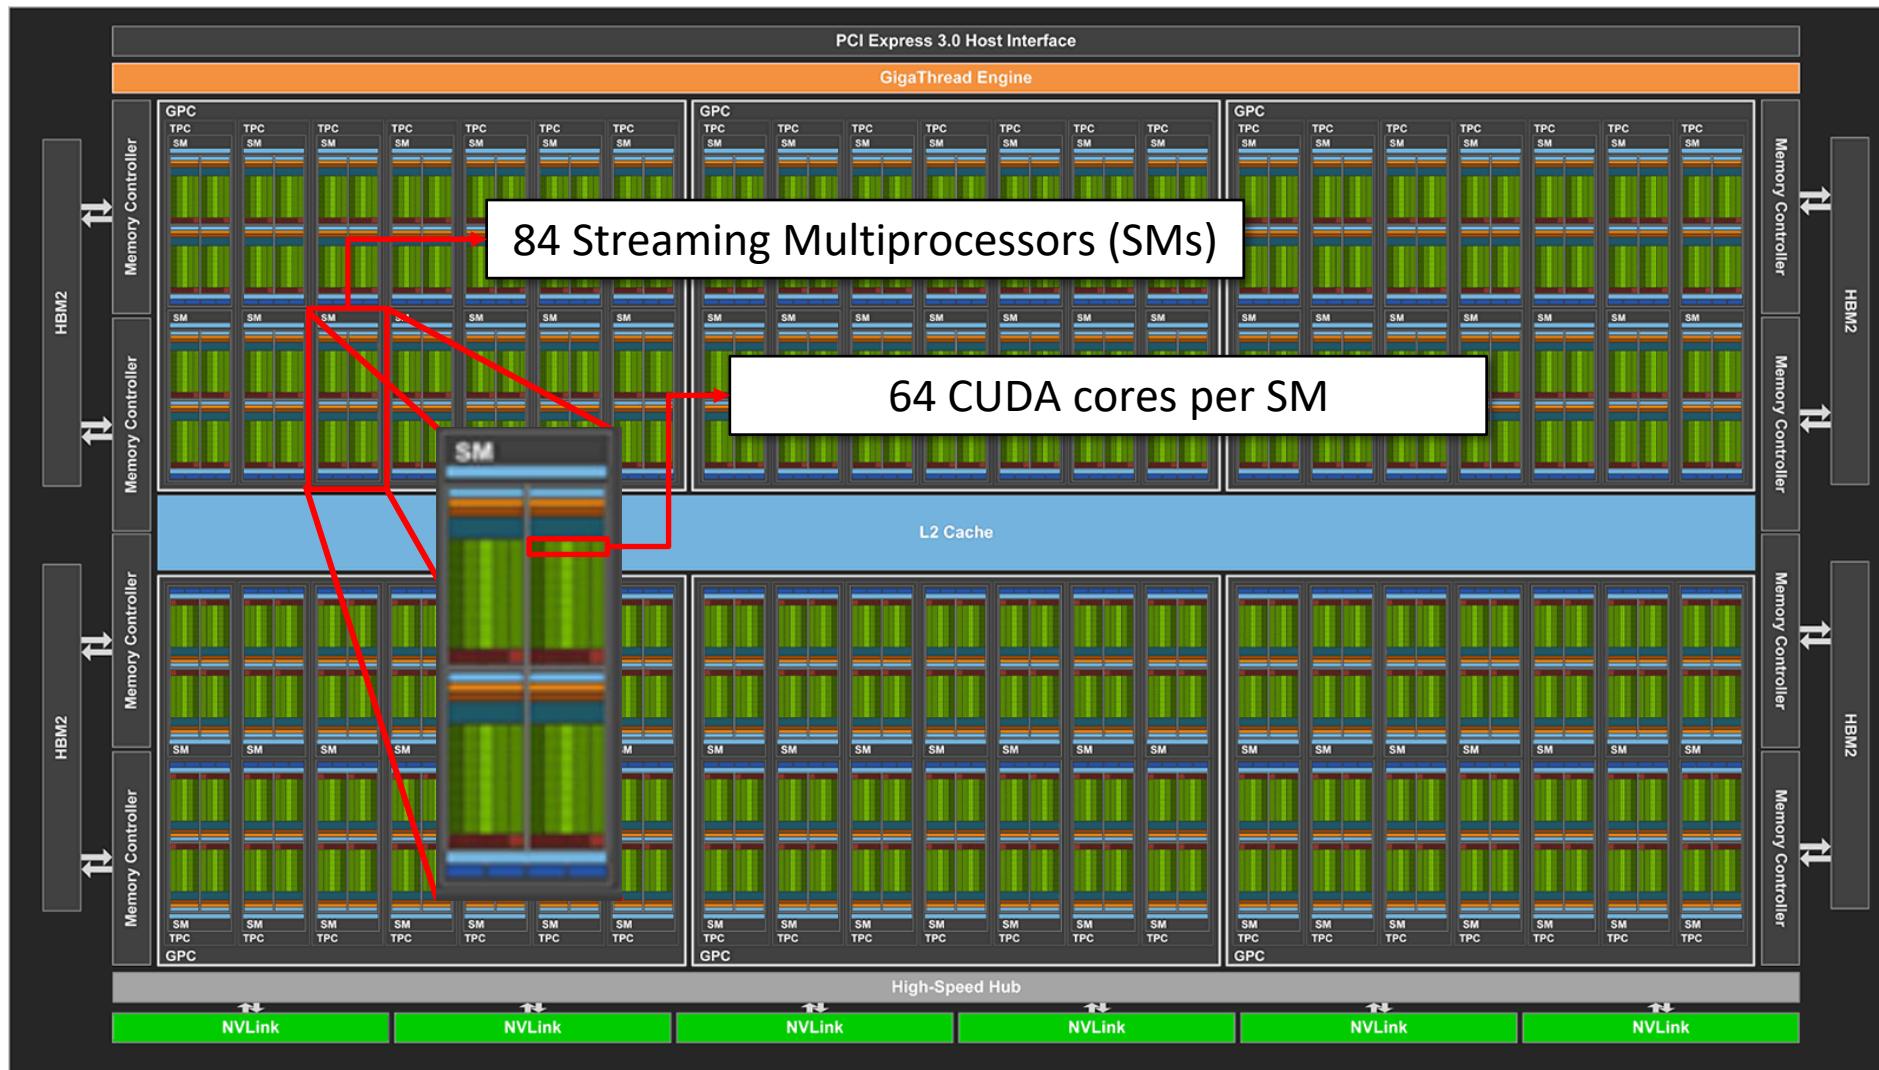
Introduction

Volta (2017)

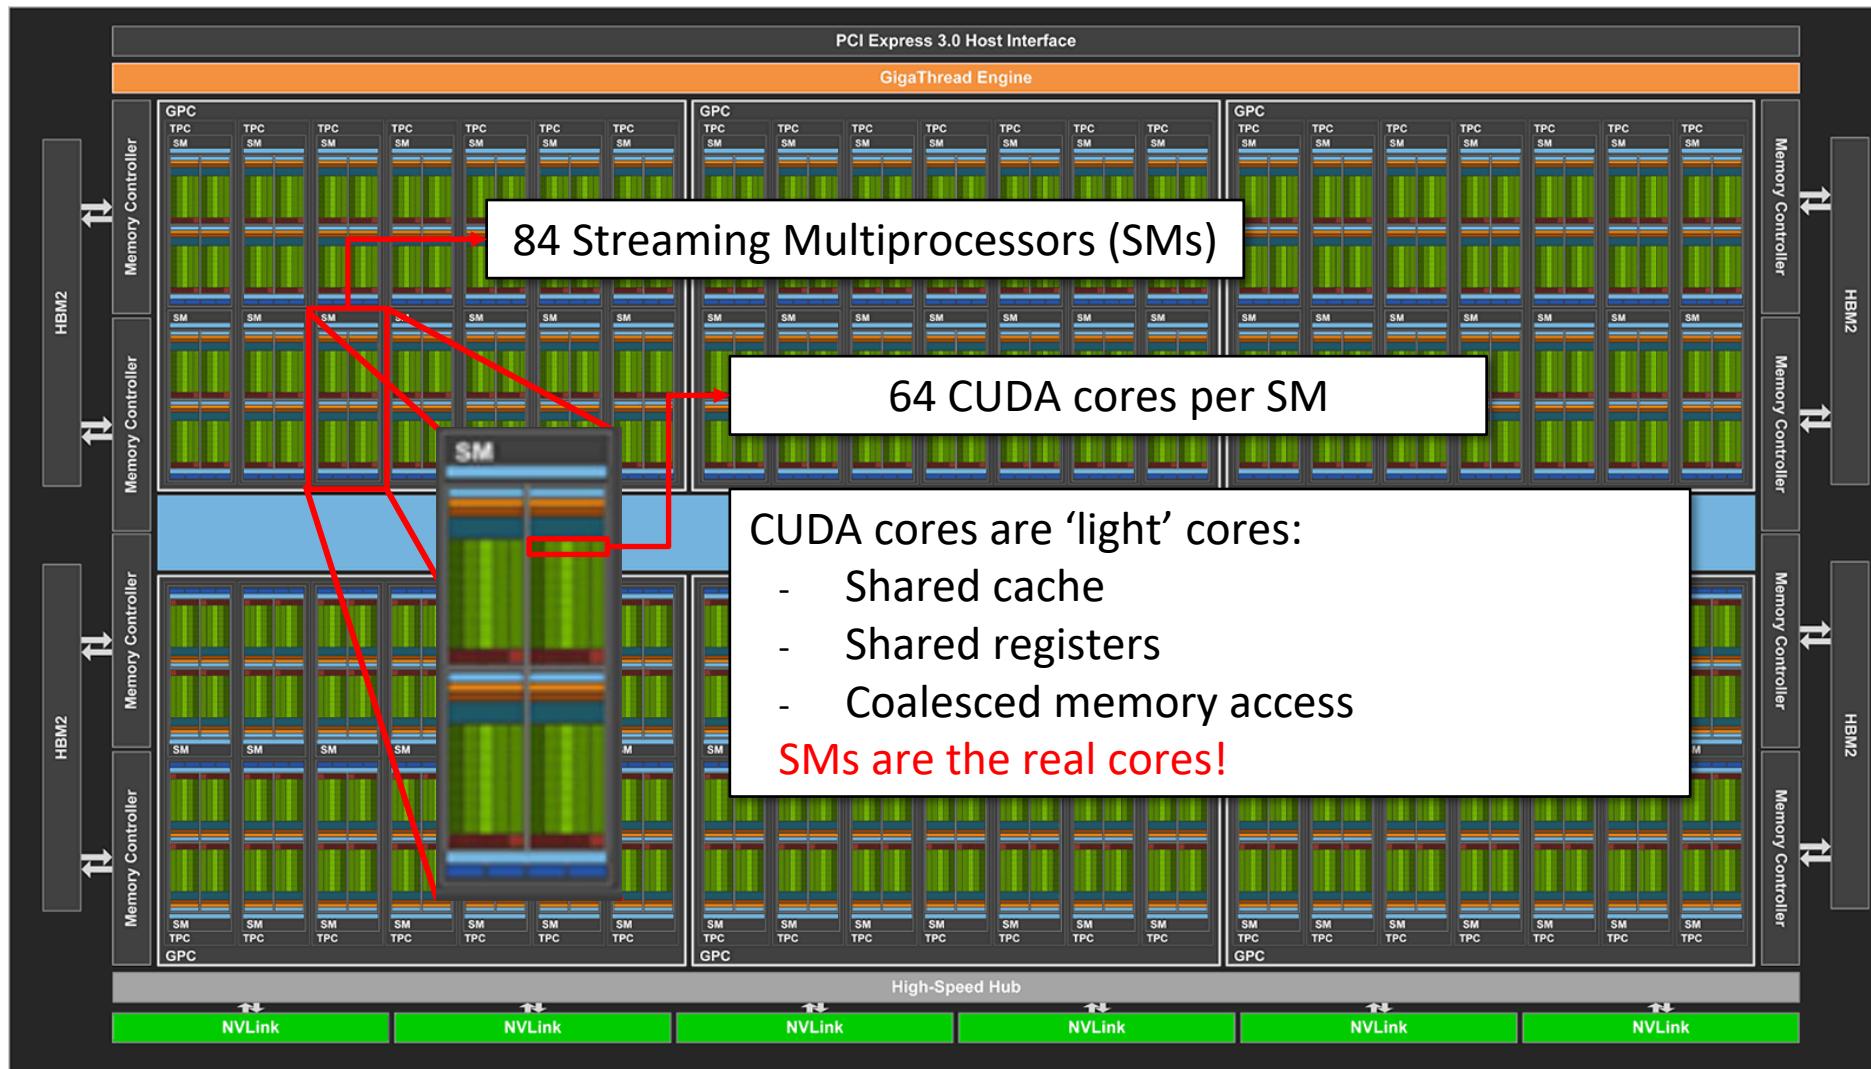


Introduction

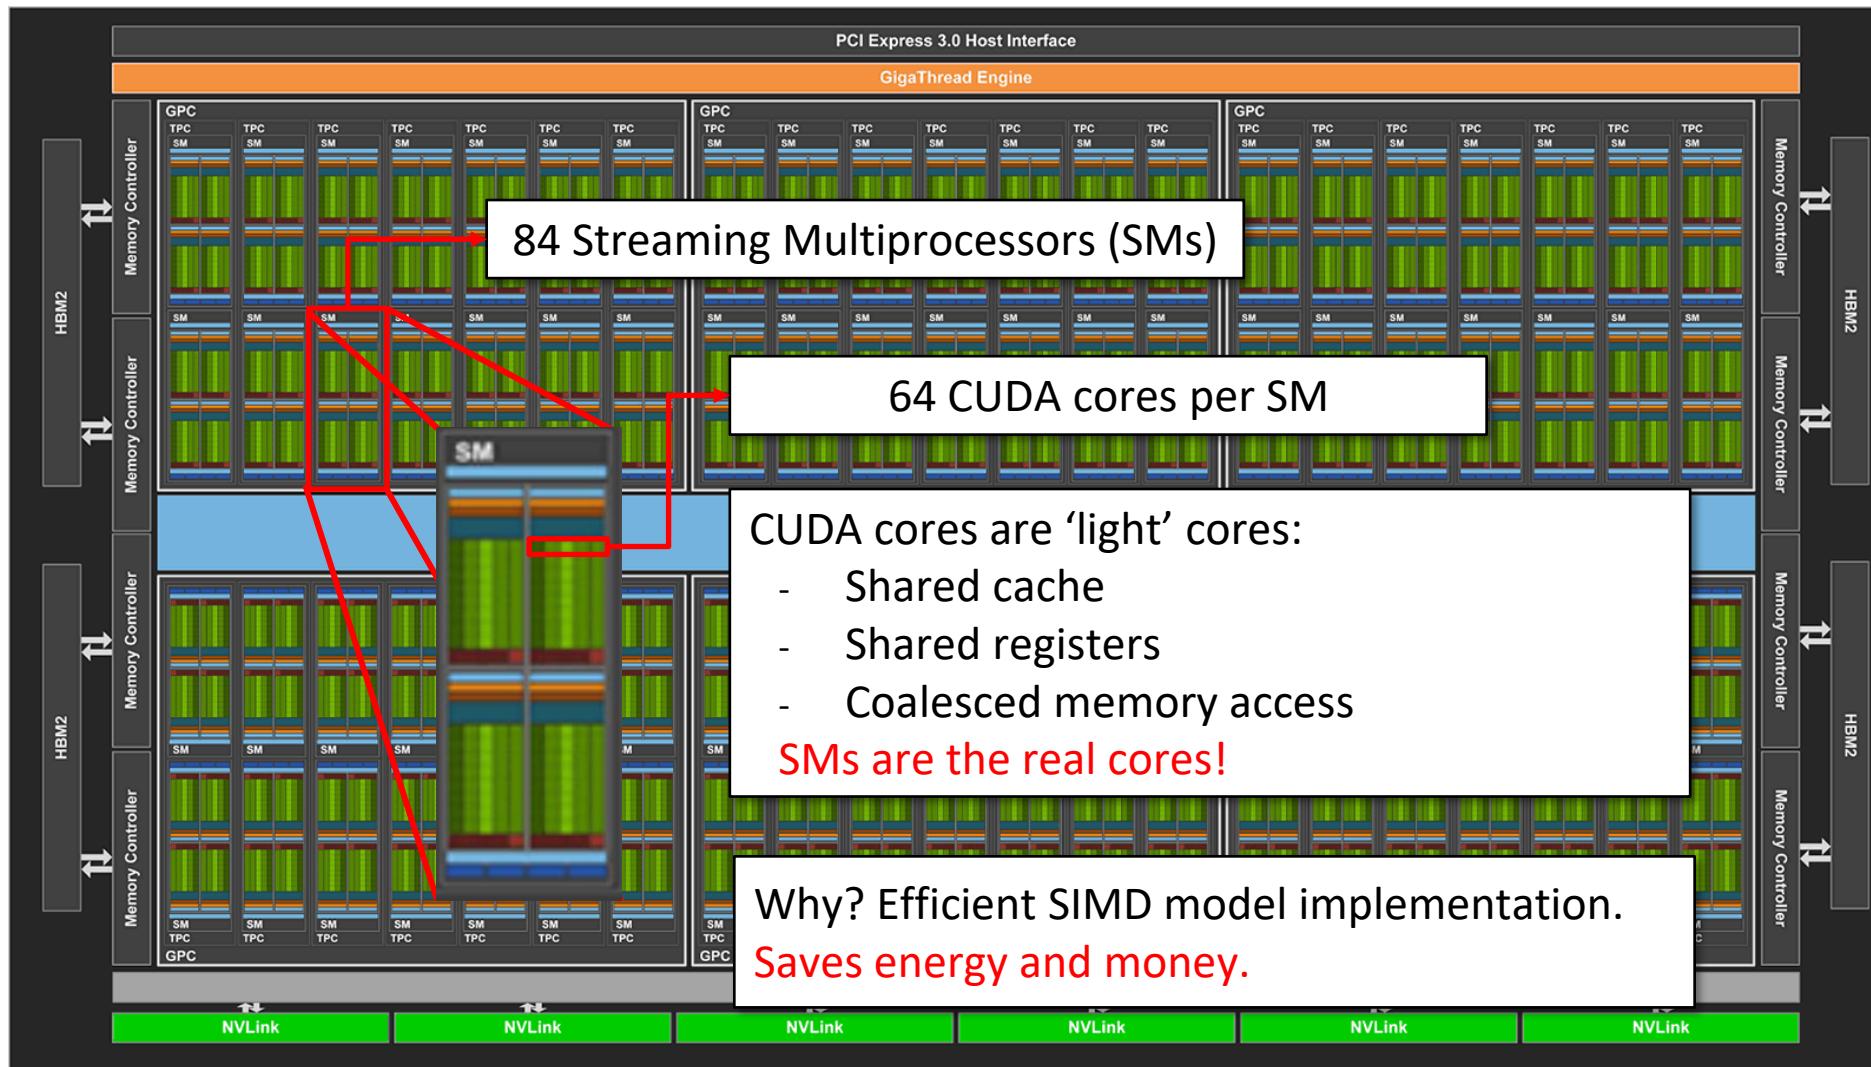
Volta (2017)



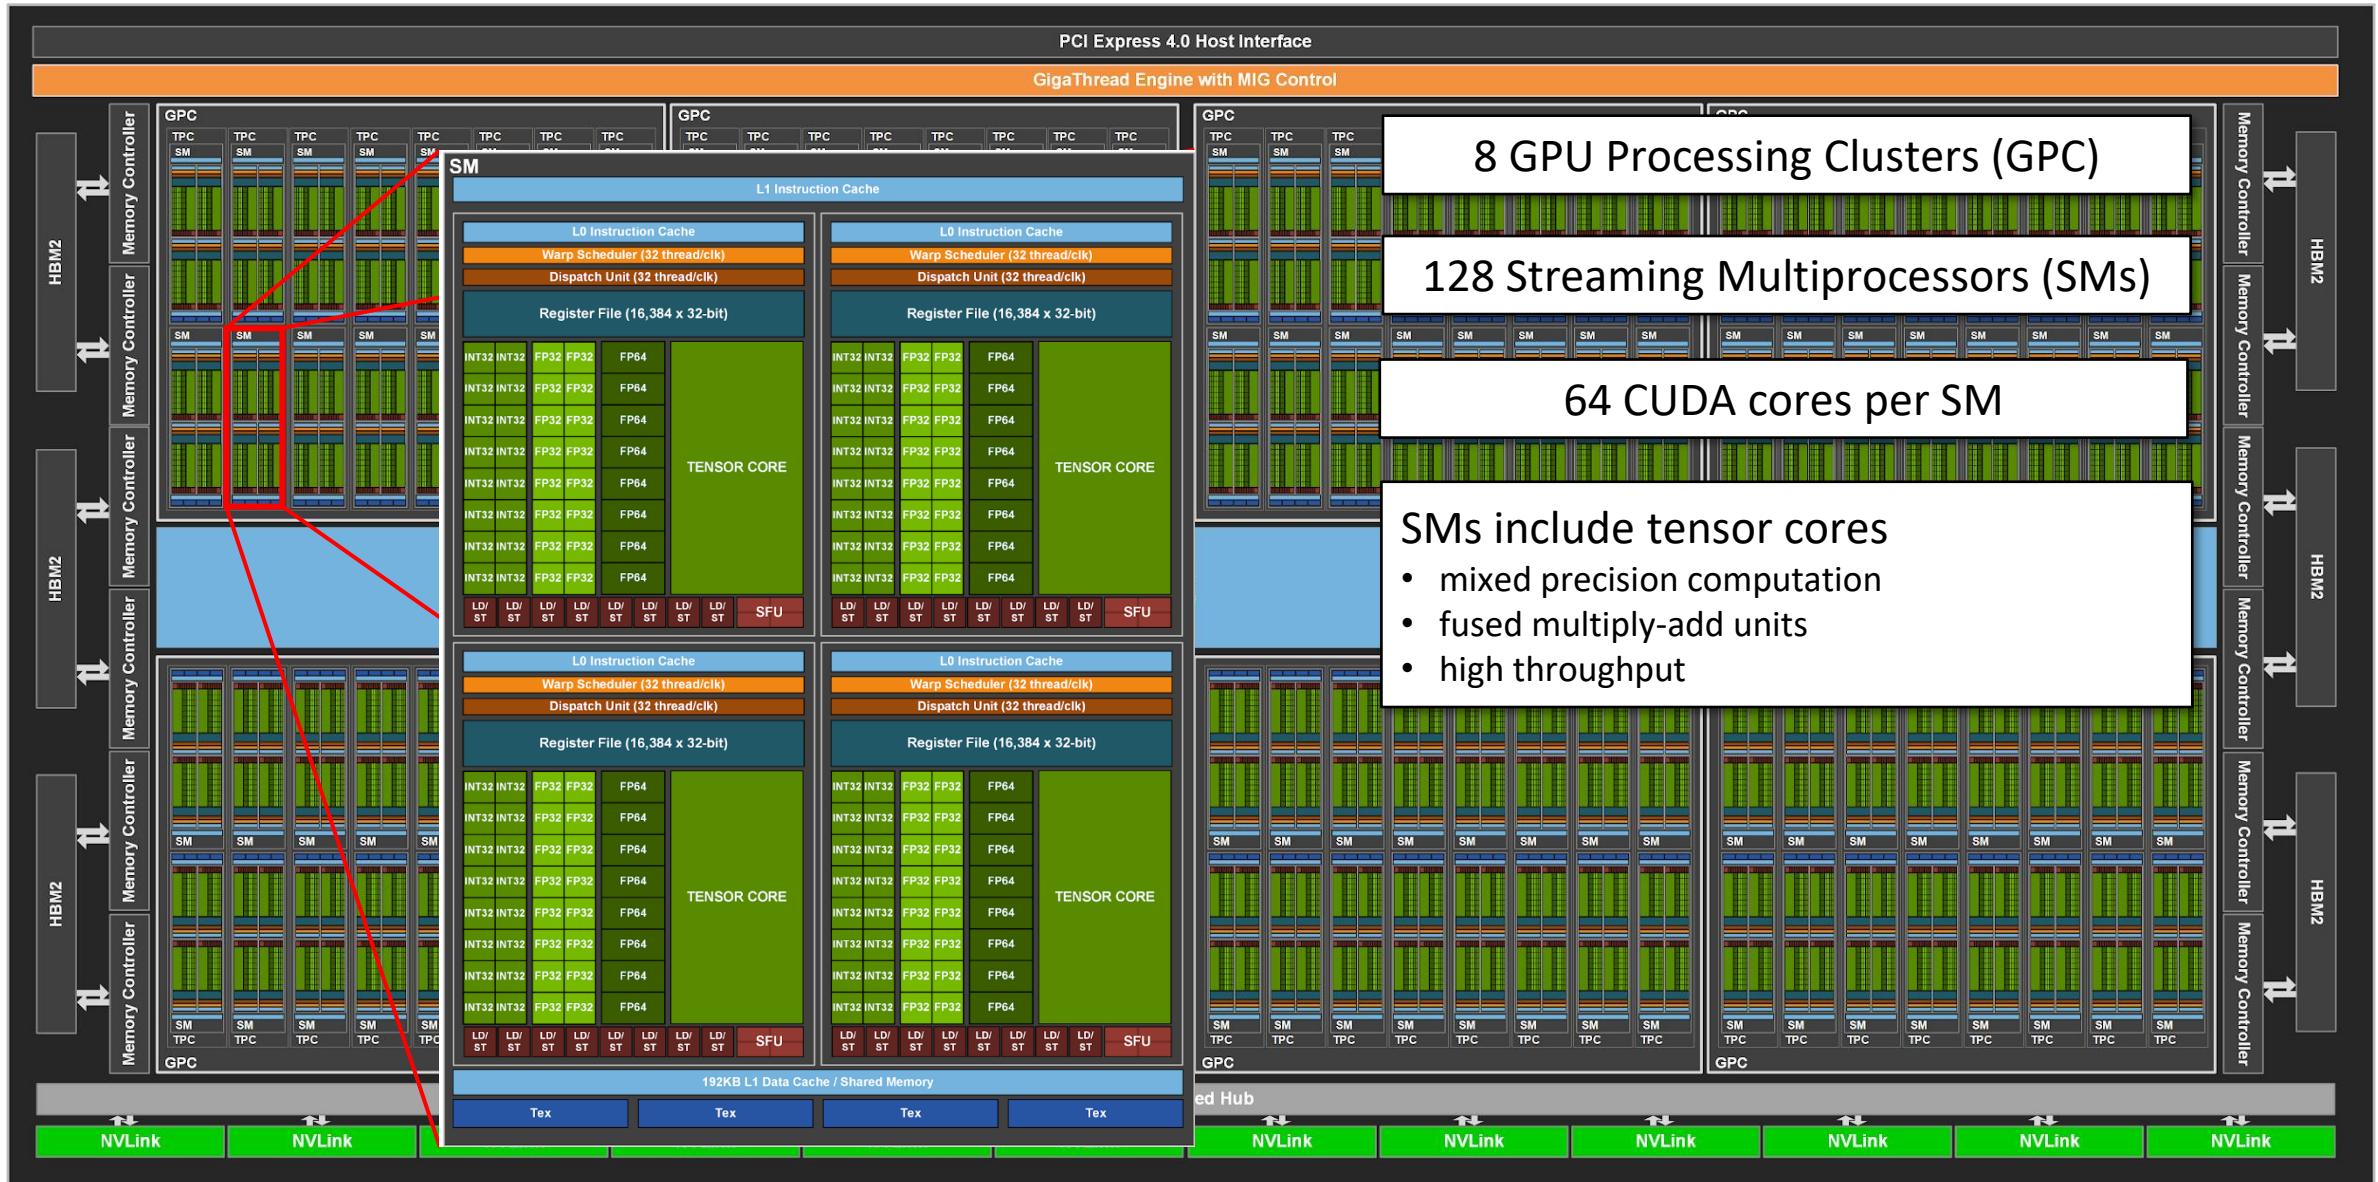
Volta (2017)



Volta (2017)

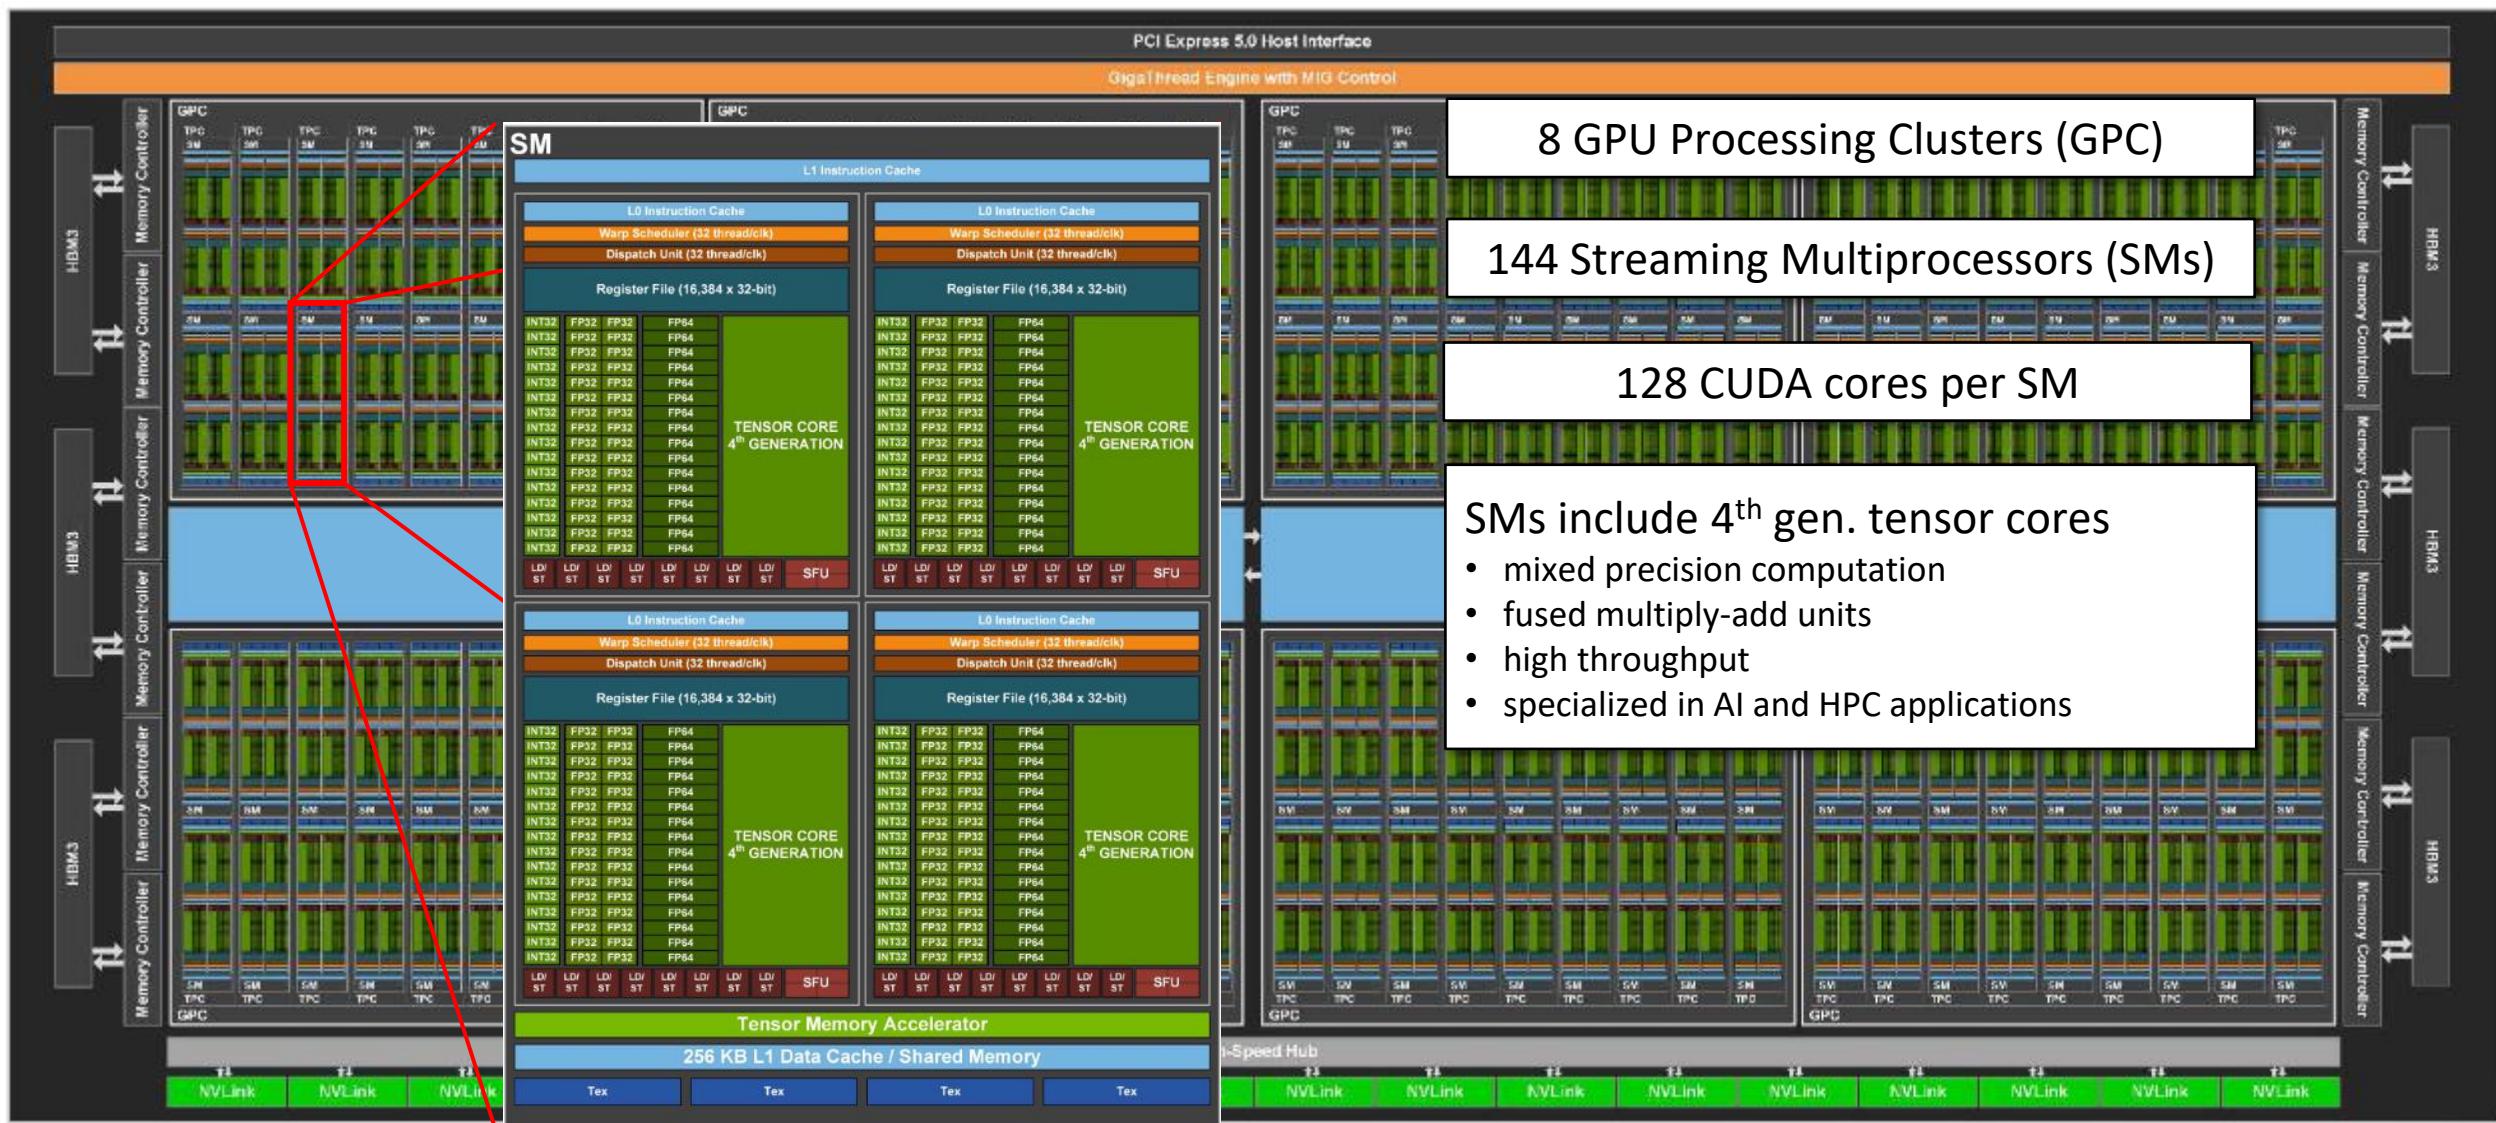


Ampere (2020)



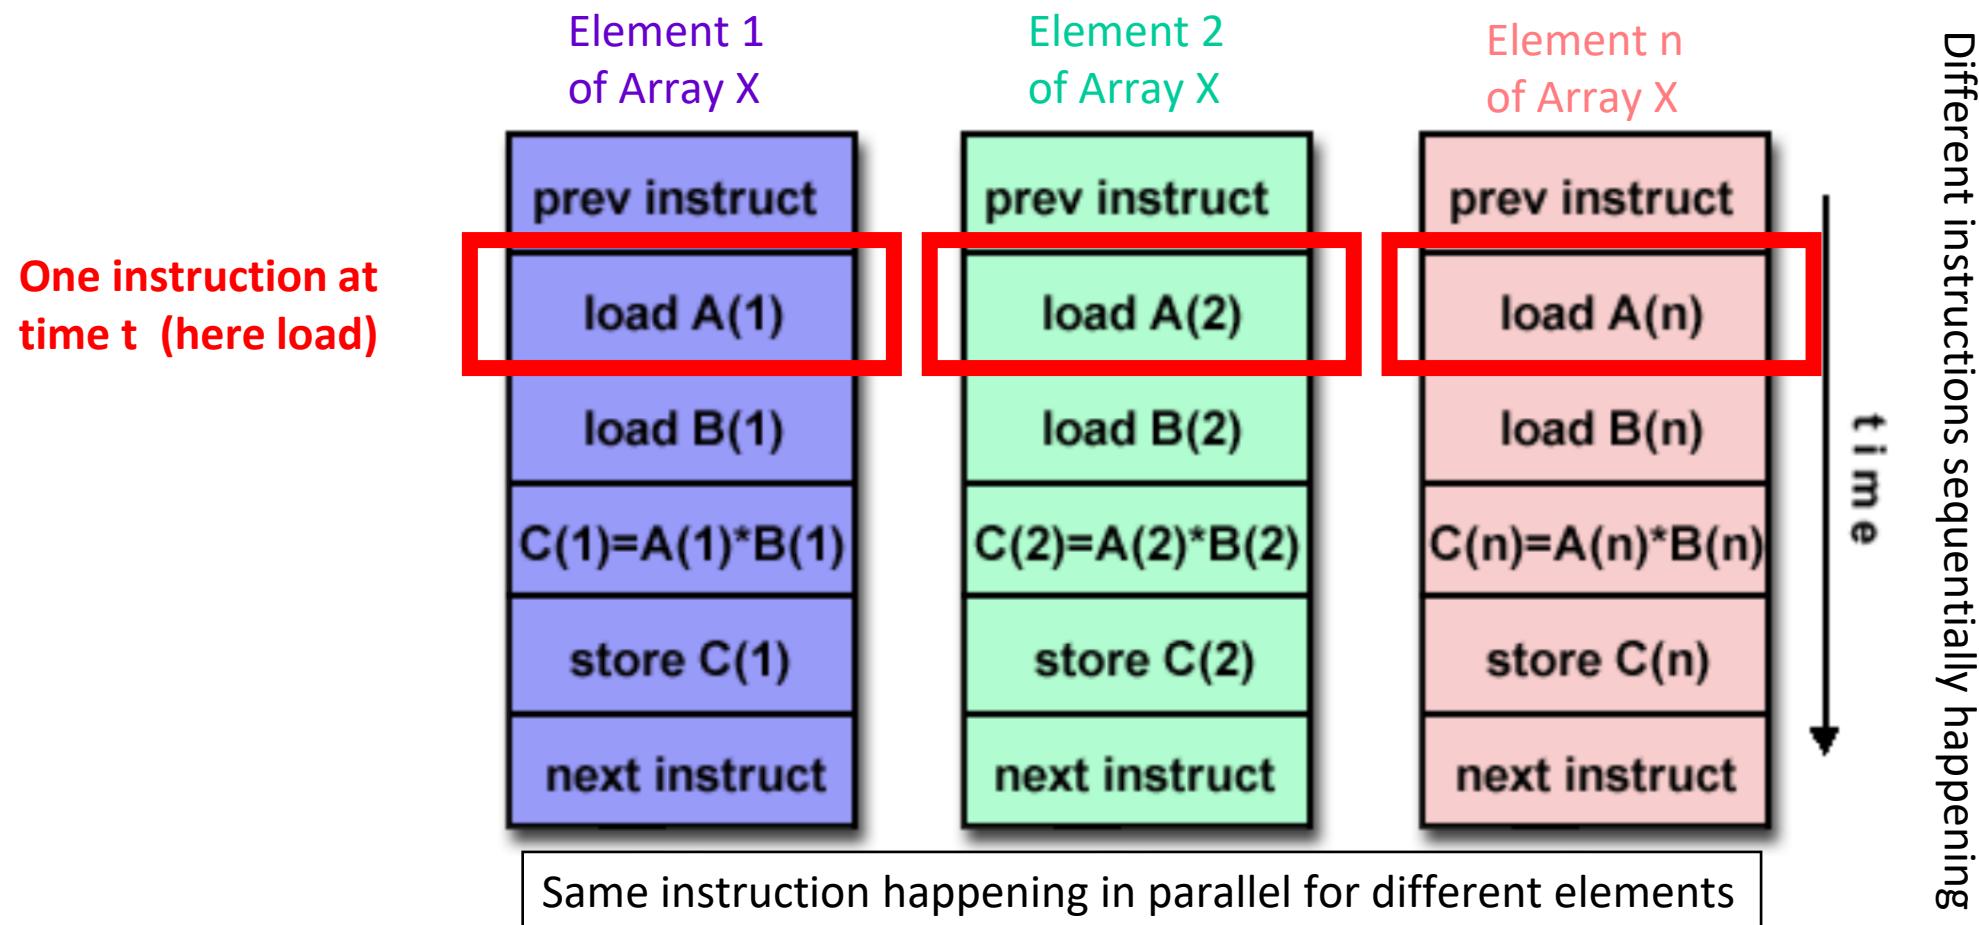
Introduction

Hopper (2022)

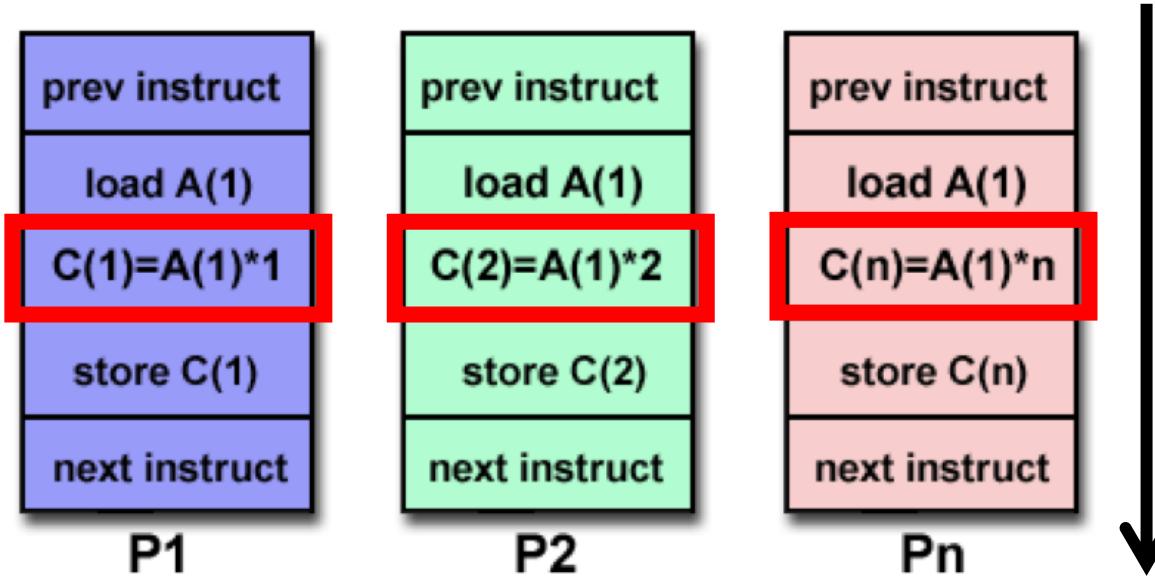


GPUs are SIMD oriented architectures

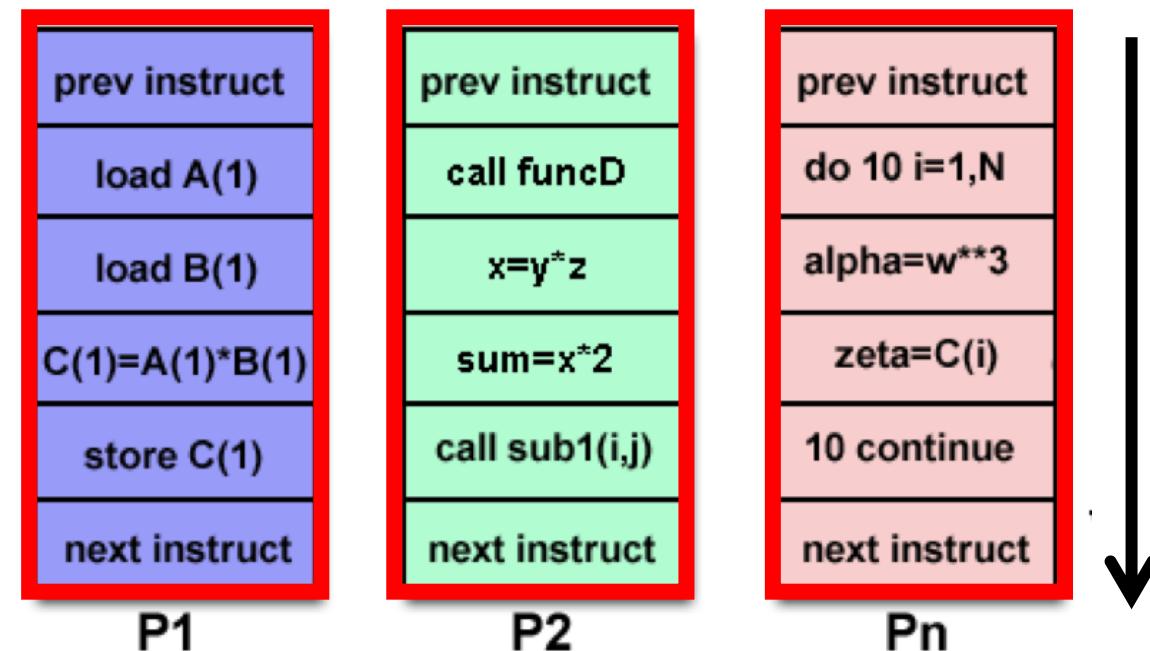
SIMD : Single Instruction Multiple Data



MISD : Multiple Instructions Single Data



MIMD : Multiple Instructions Multiple Data



GPUs are **NOT** MISD/MIMD oriented architectures

CPU can do MIMD & SIMD

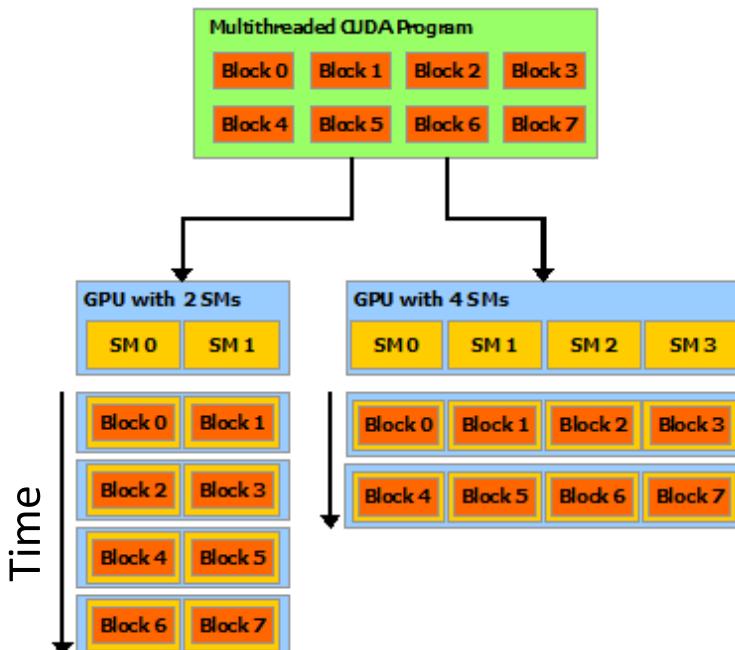
GPU can do SIMD only

Introduction

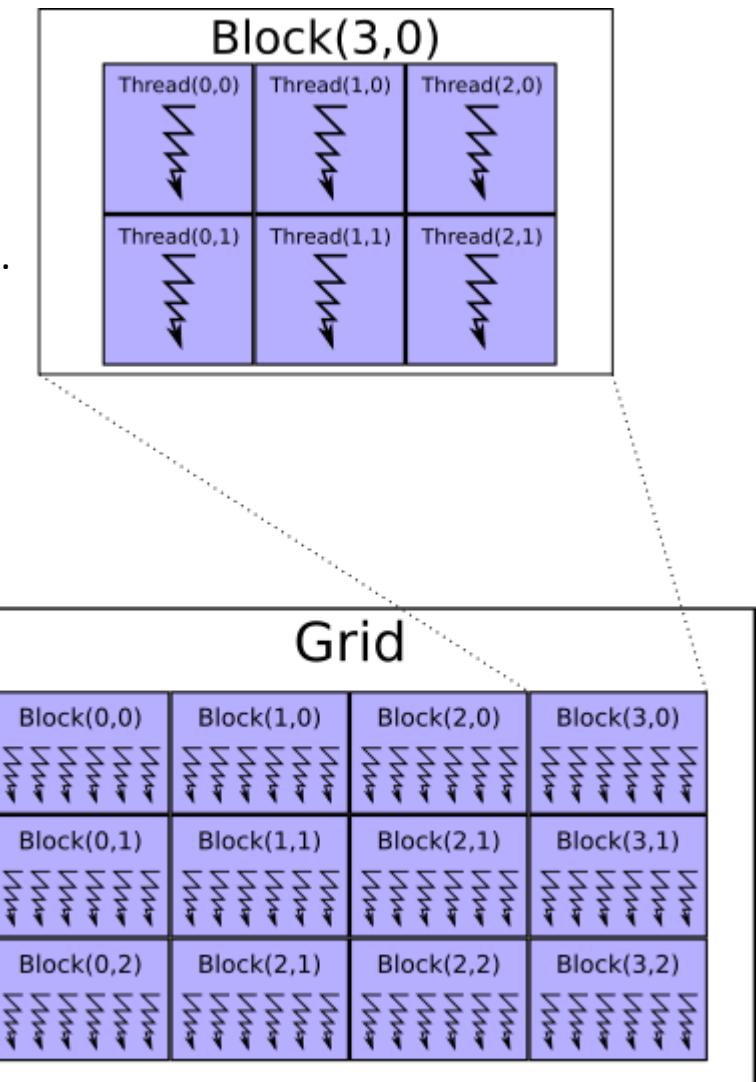
Mapping the hardware architecture into the programming architecture

As a GPU programmer you will :

- define an instruction applied on an element of an array : **thread**.
- run many threads in parallel within a shared memory environment : **blocks**.
(block = doing the same operation in parallel on a subset of elements of your array).
- manage many blocks in parallel (many blocks on the same SM, or many SM), or sequentially if your array exceed your GPU capacity.



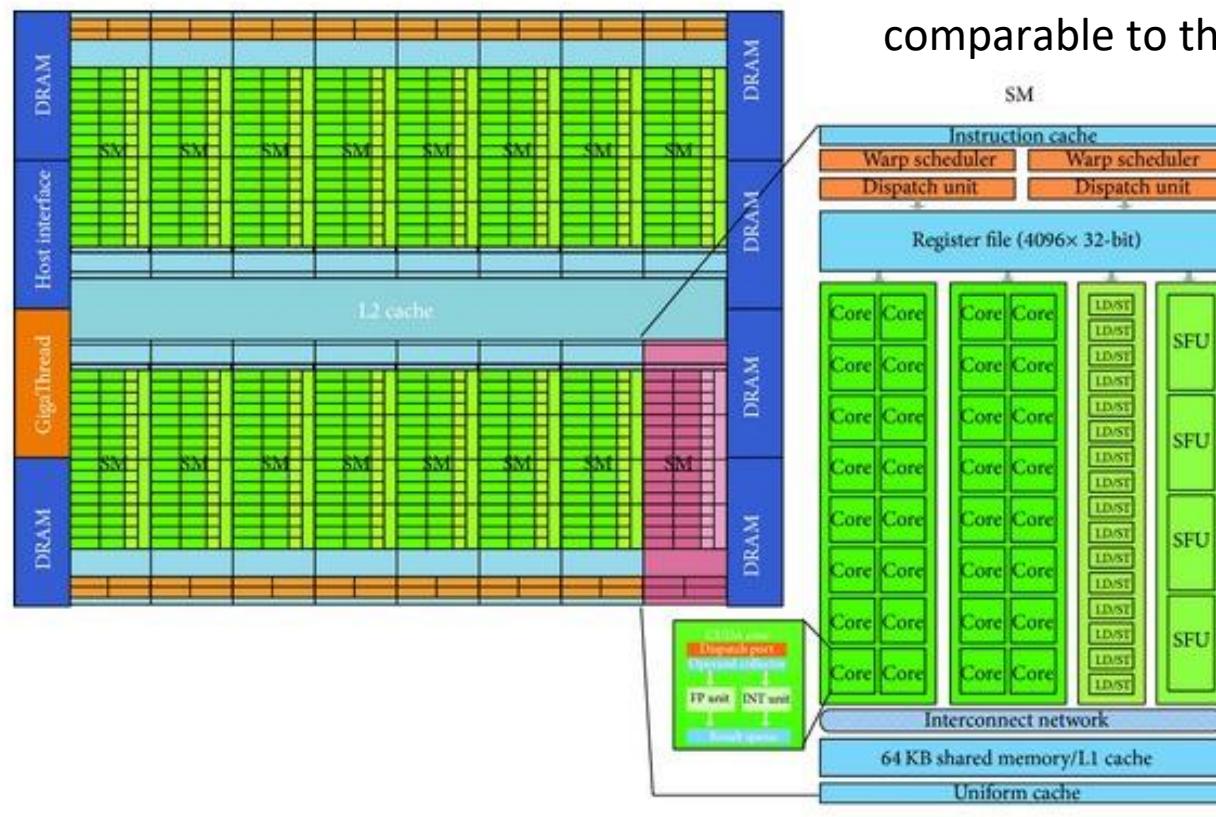
Warp: Minimum amount of threads (32) that will be executed simultaneously per SM



Introduction

How does it look like on a real architecture?

Nvidia Fermi GPU (2011)



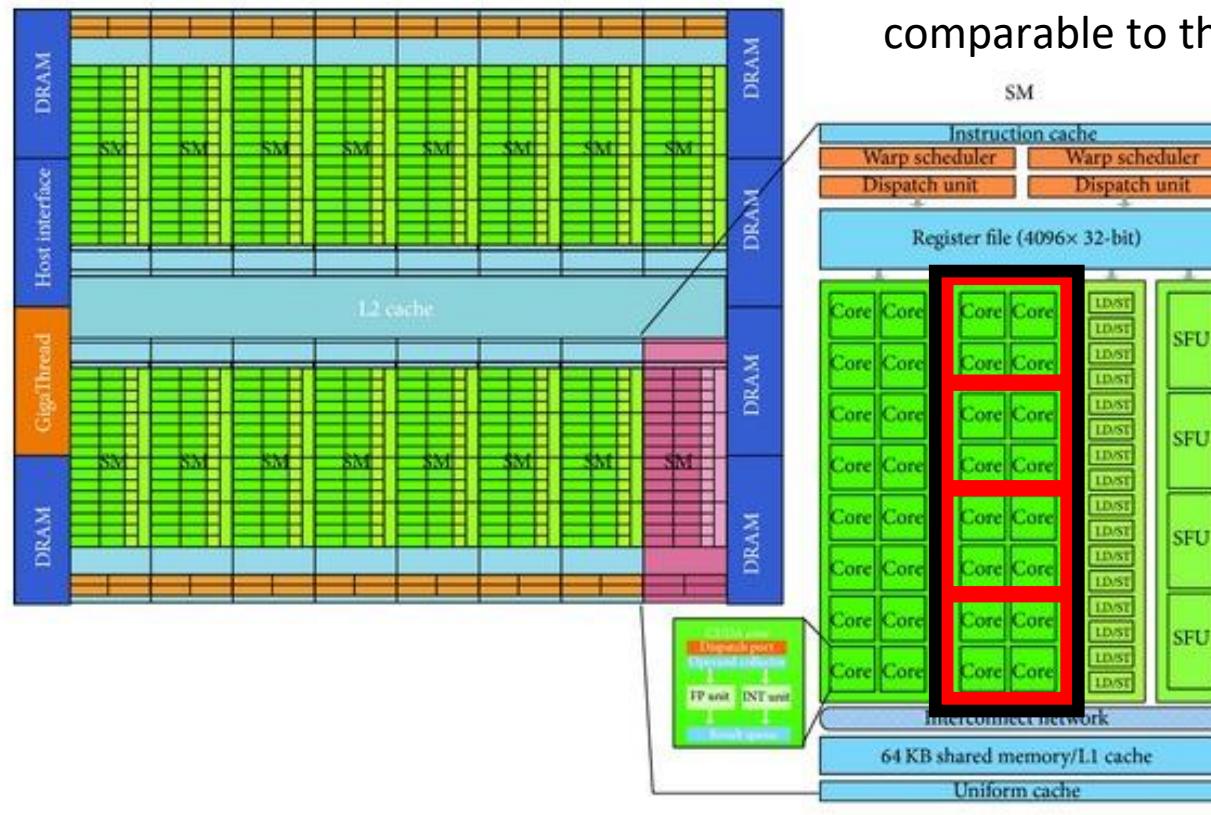
Here it is not a core, it is a "CUDA core" so not directly comparable to the CPU core

Source: 3D Data Denoising via Nonlocal Means Filter by Using Parallel GPU Strategies
DOI: 10.1155/2014/523862

Introduction

How does it look like on a real architecture?

Nvidia Fermi GPU (2011)



Here it is not a core, it is a "CUDA core" so not directly comparable to the CPU core

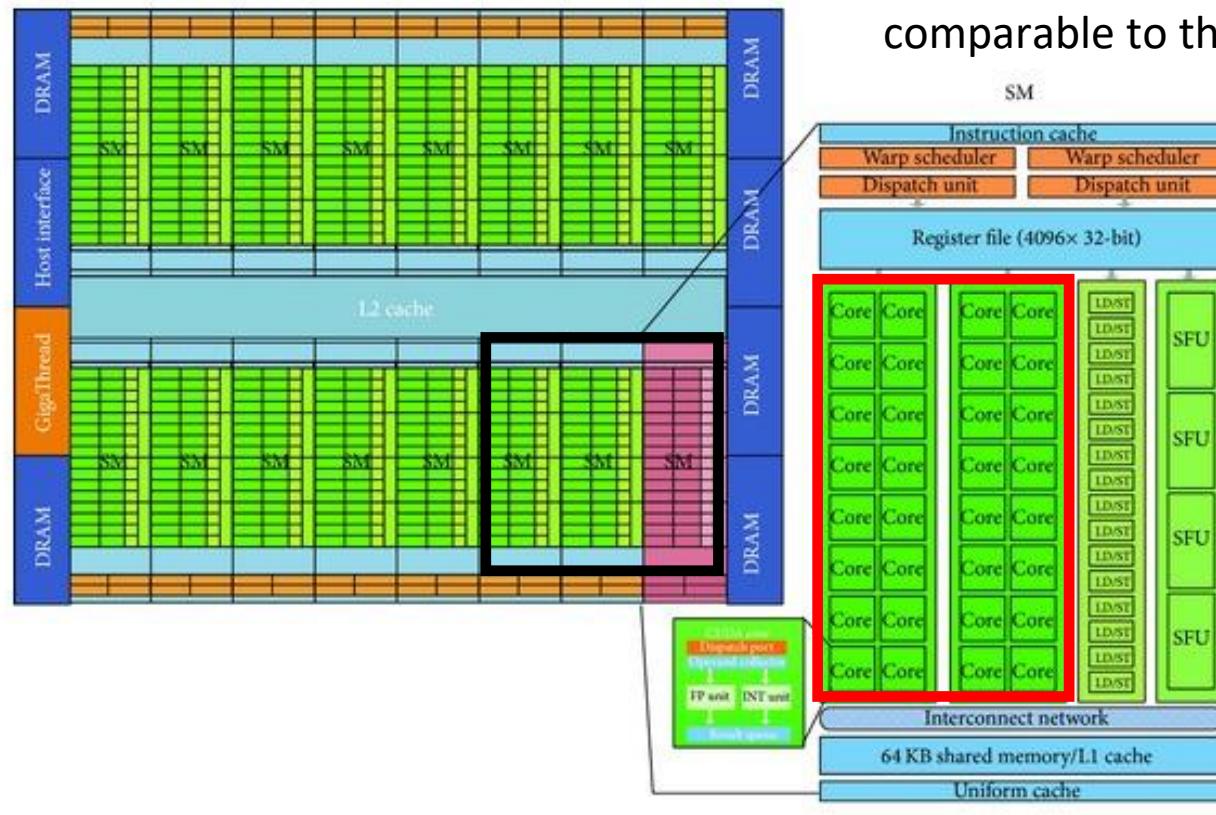
Grid

Block

Introduction

How does it look like on a real architecture?

Nvidia Fermi GPU (2011)



Here it is not a core, it is a "CUDA core" so not directly comparable to the CPU core

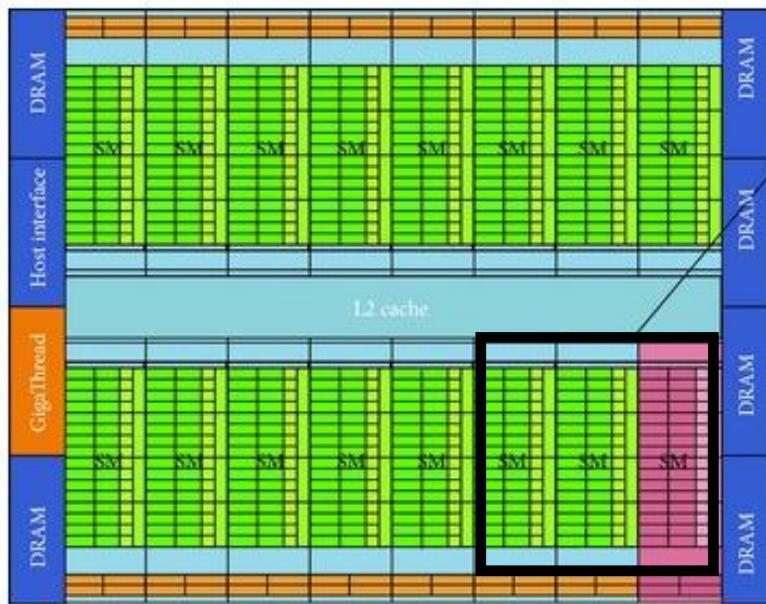
Grid

Block

Introduction

How does it look like on a real architecture?

Nvidia Fermi GPU (2011)



A grid can encompass many SM but a block can't use more CUDA-cores than those in a SM since a block is a **shared memory space**.

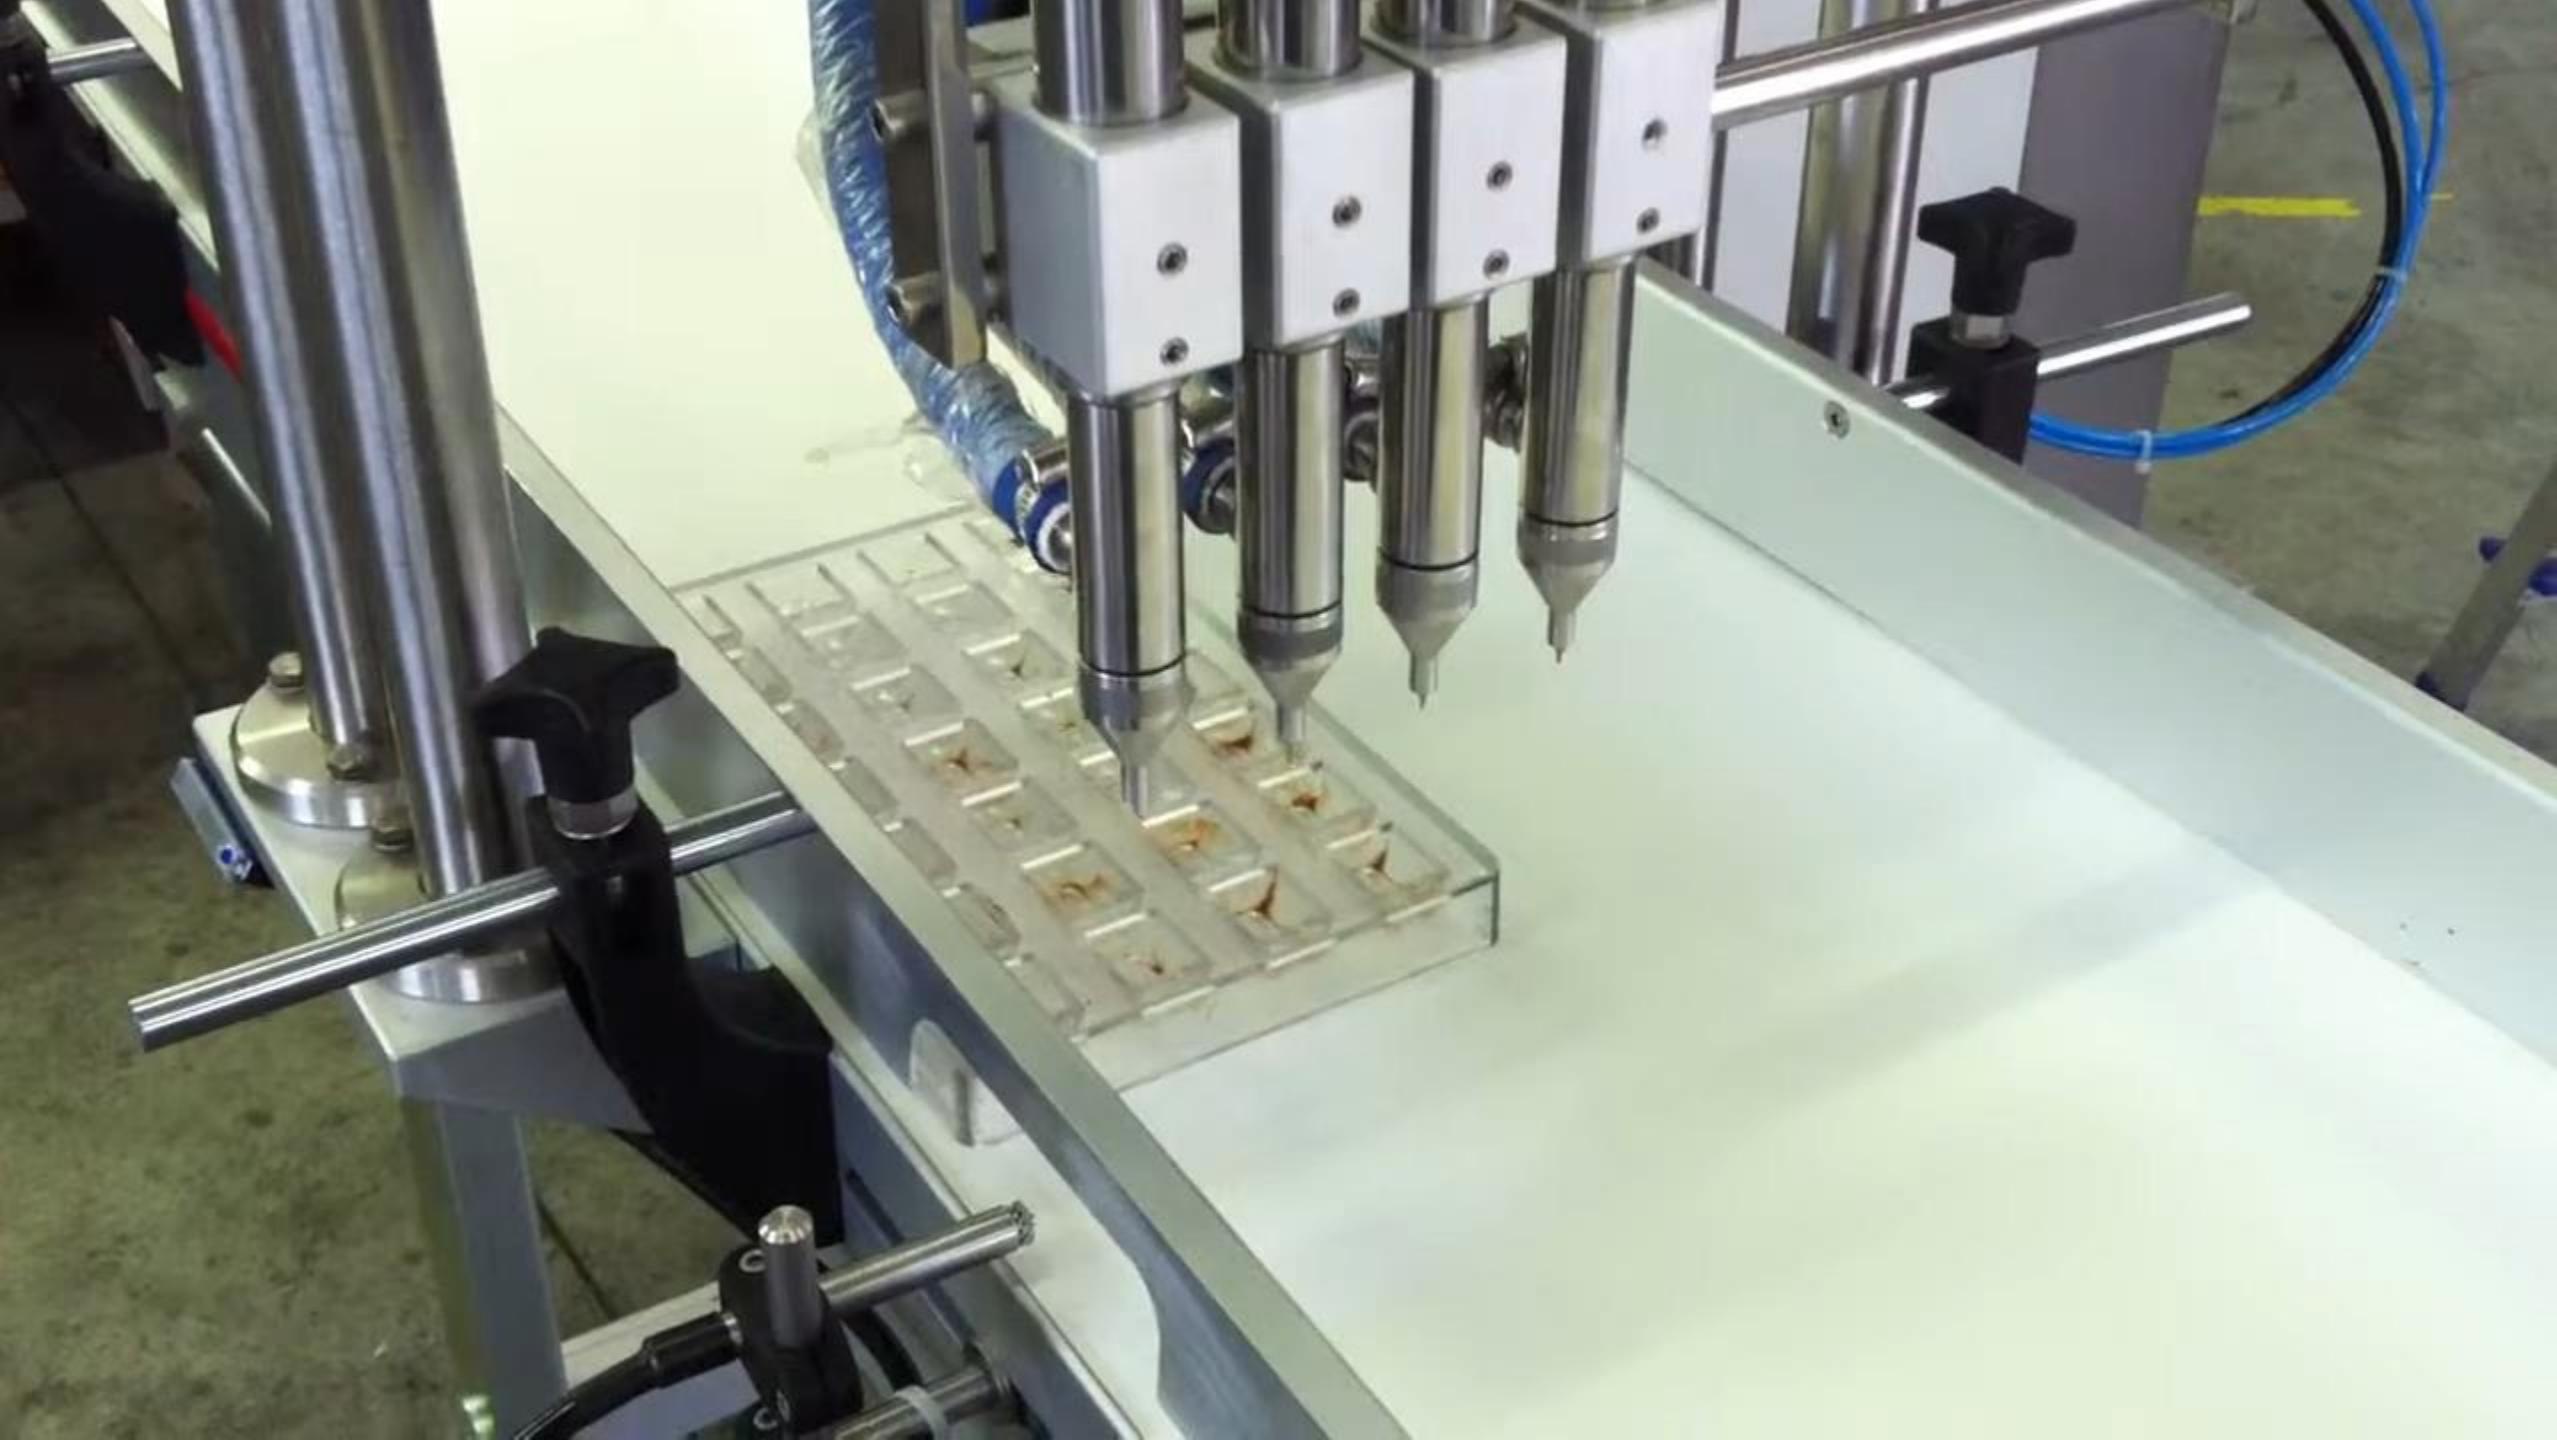
Here it is not a core, it is a "CUDA core" so not directly comparable to the CPU core



Grid

Block

You can also set the number of threads per block, but this number is not easily linkable to the number of CUDA cores.
Yet a **block size (number of threads per block)** should be a multiple of 32



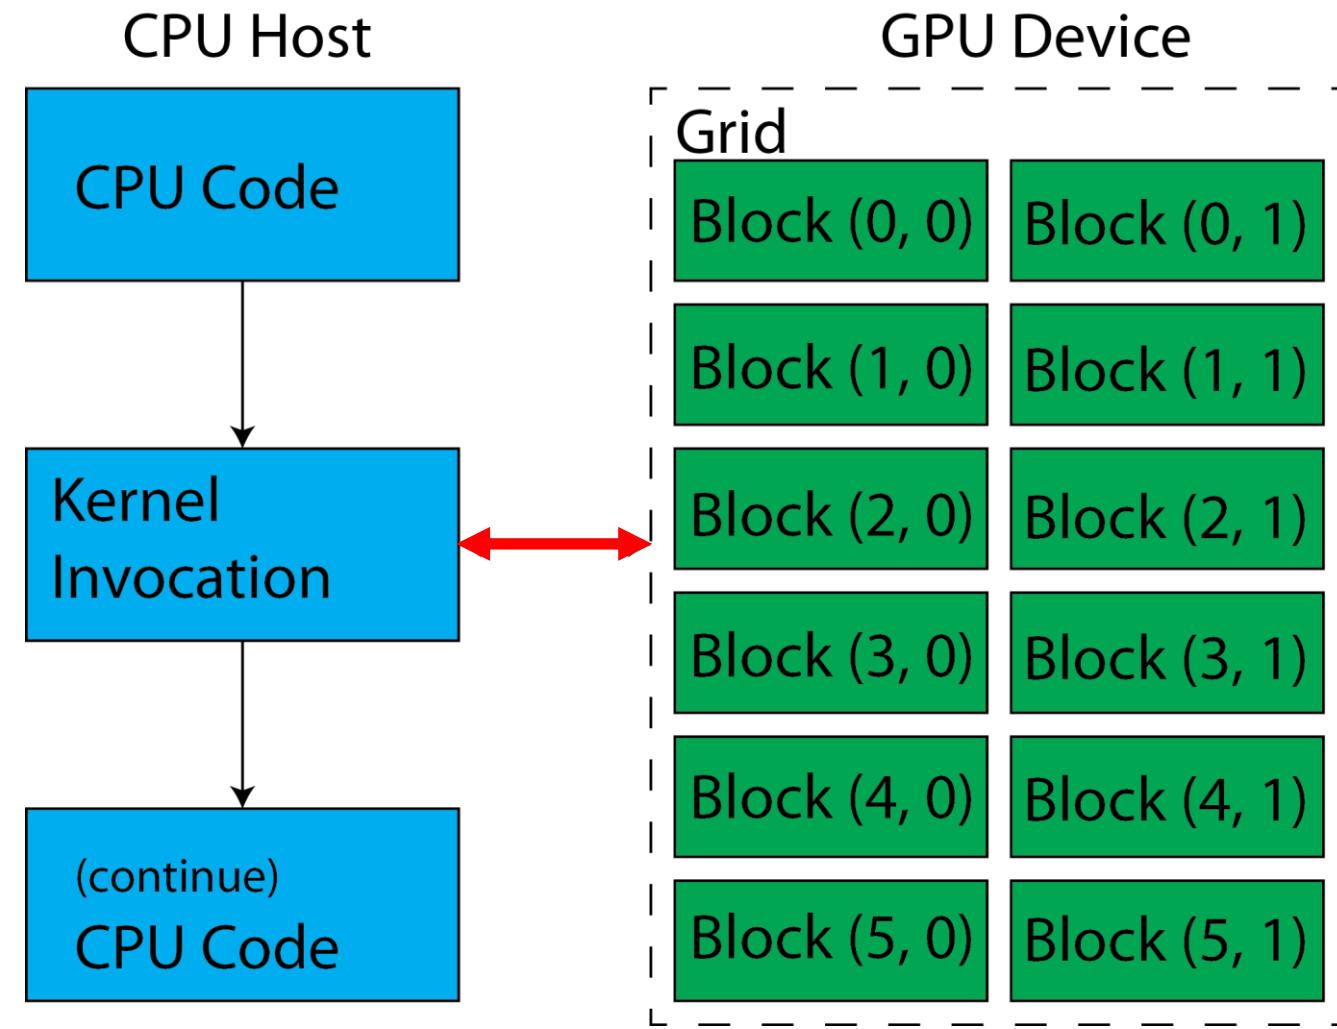
Outline

Day 1

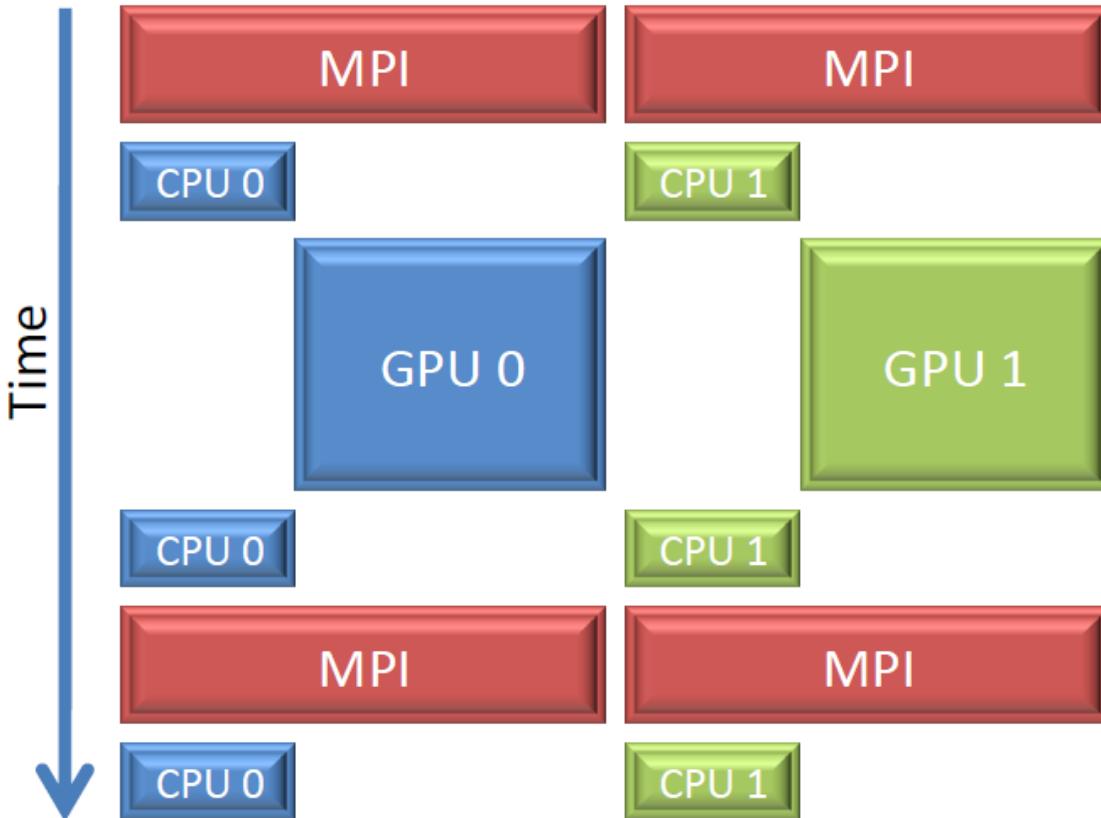
(09:00 – 13:00)

- ❑ Introduction
- ❑ Coding a CUDA kernel
- ❑ Example
 - ❑ Matrix Multiplication + Exercise
- ❑ Real world implementations
 - ❑ Genomes simulation
 - ❑ Random numbers generation
- ❑ Homework

A mapping at the core of CPU/GPU communication: the **kernel**

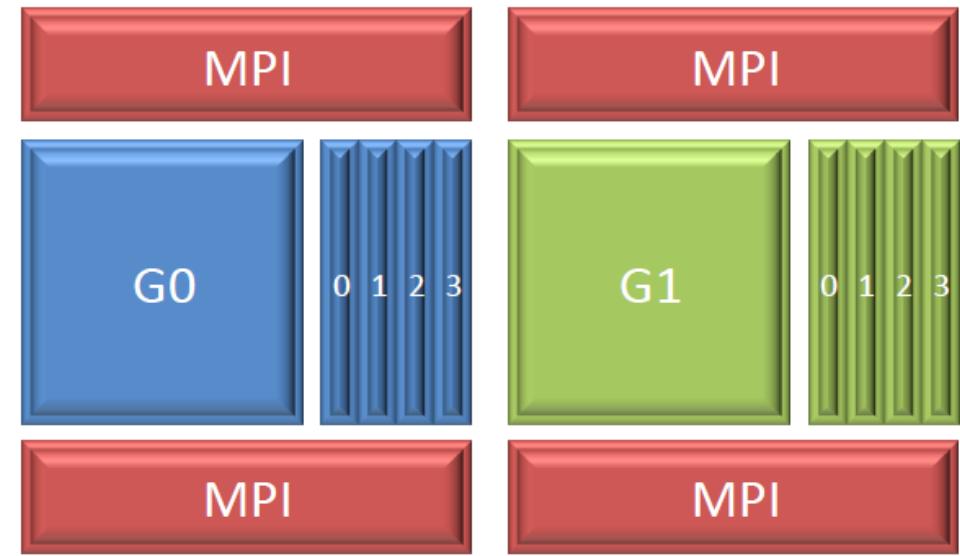


So-So Hybridization



- Neglects CPU
- Suffers from Amdahl's Law

Better Hybridization



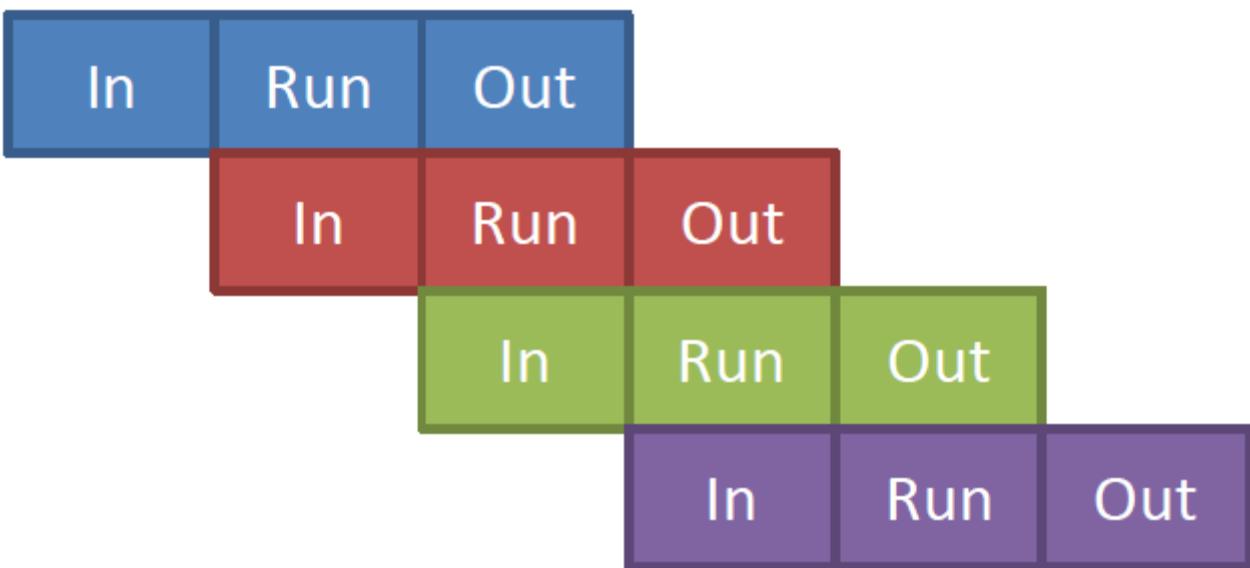
- Overlap CPU/GPU work and data movement
- Even better if you can overlap MPI comms

Most GPU operations are asynchronous from the CPU code
(i.e. the CPU can be busy doing other things!)

Synchronous execution (1 stream)

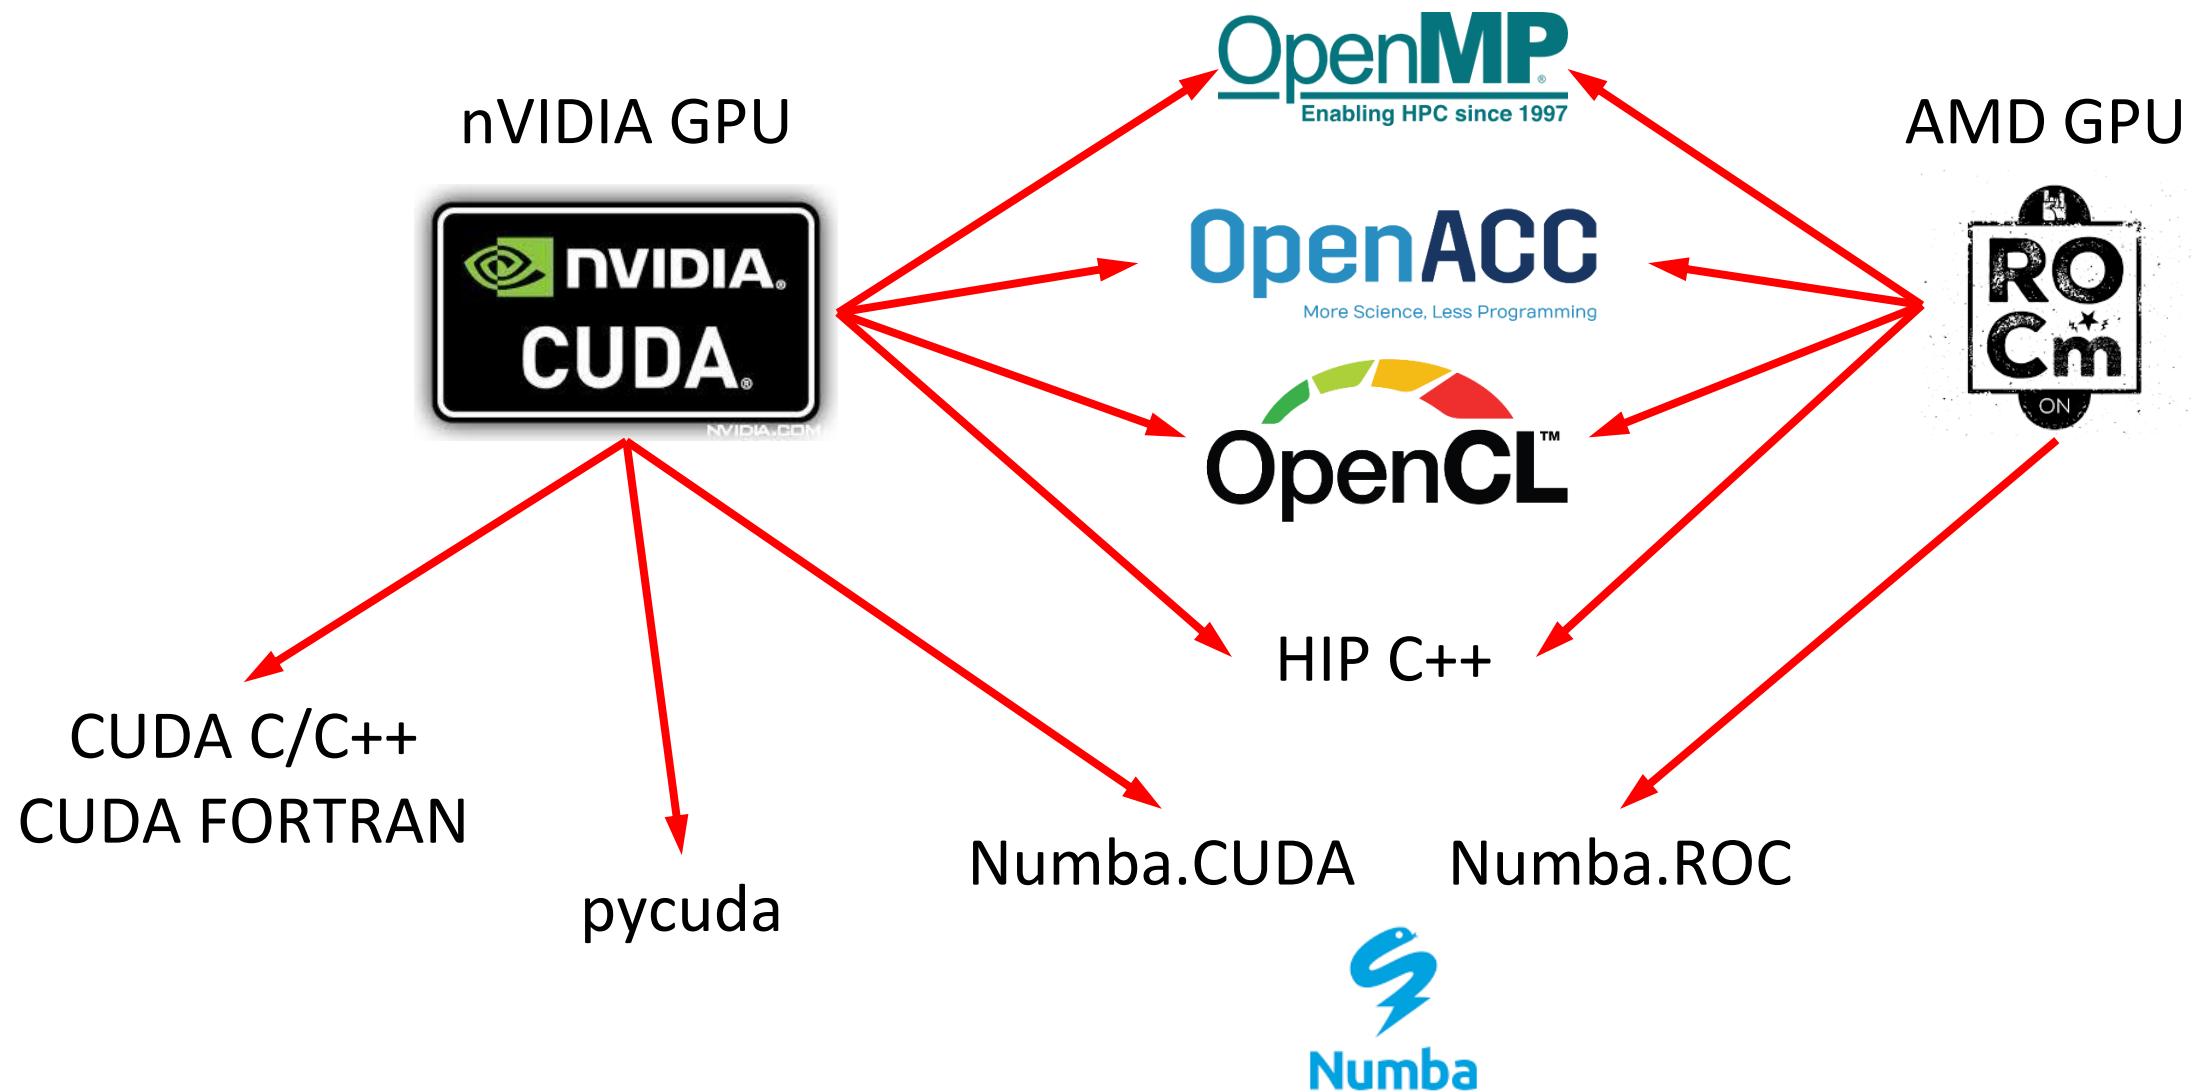


Asynchronous execution (3 streams)

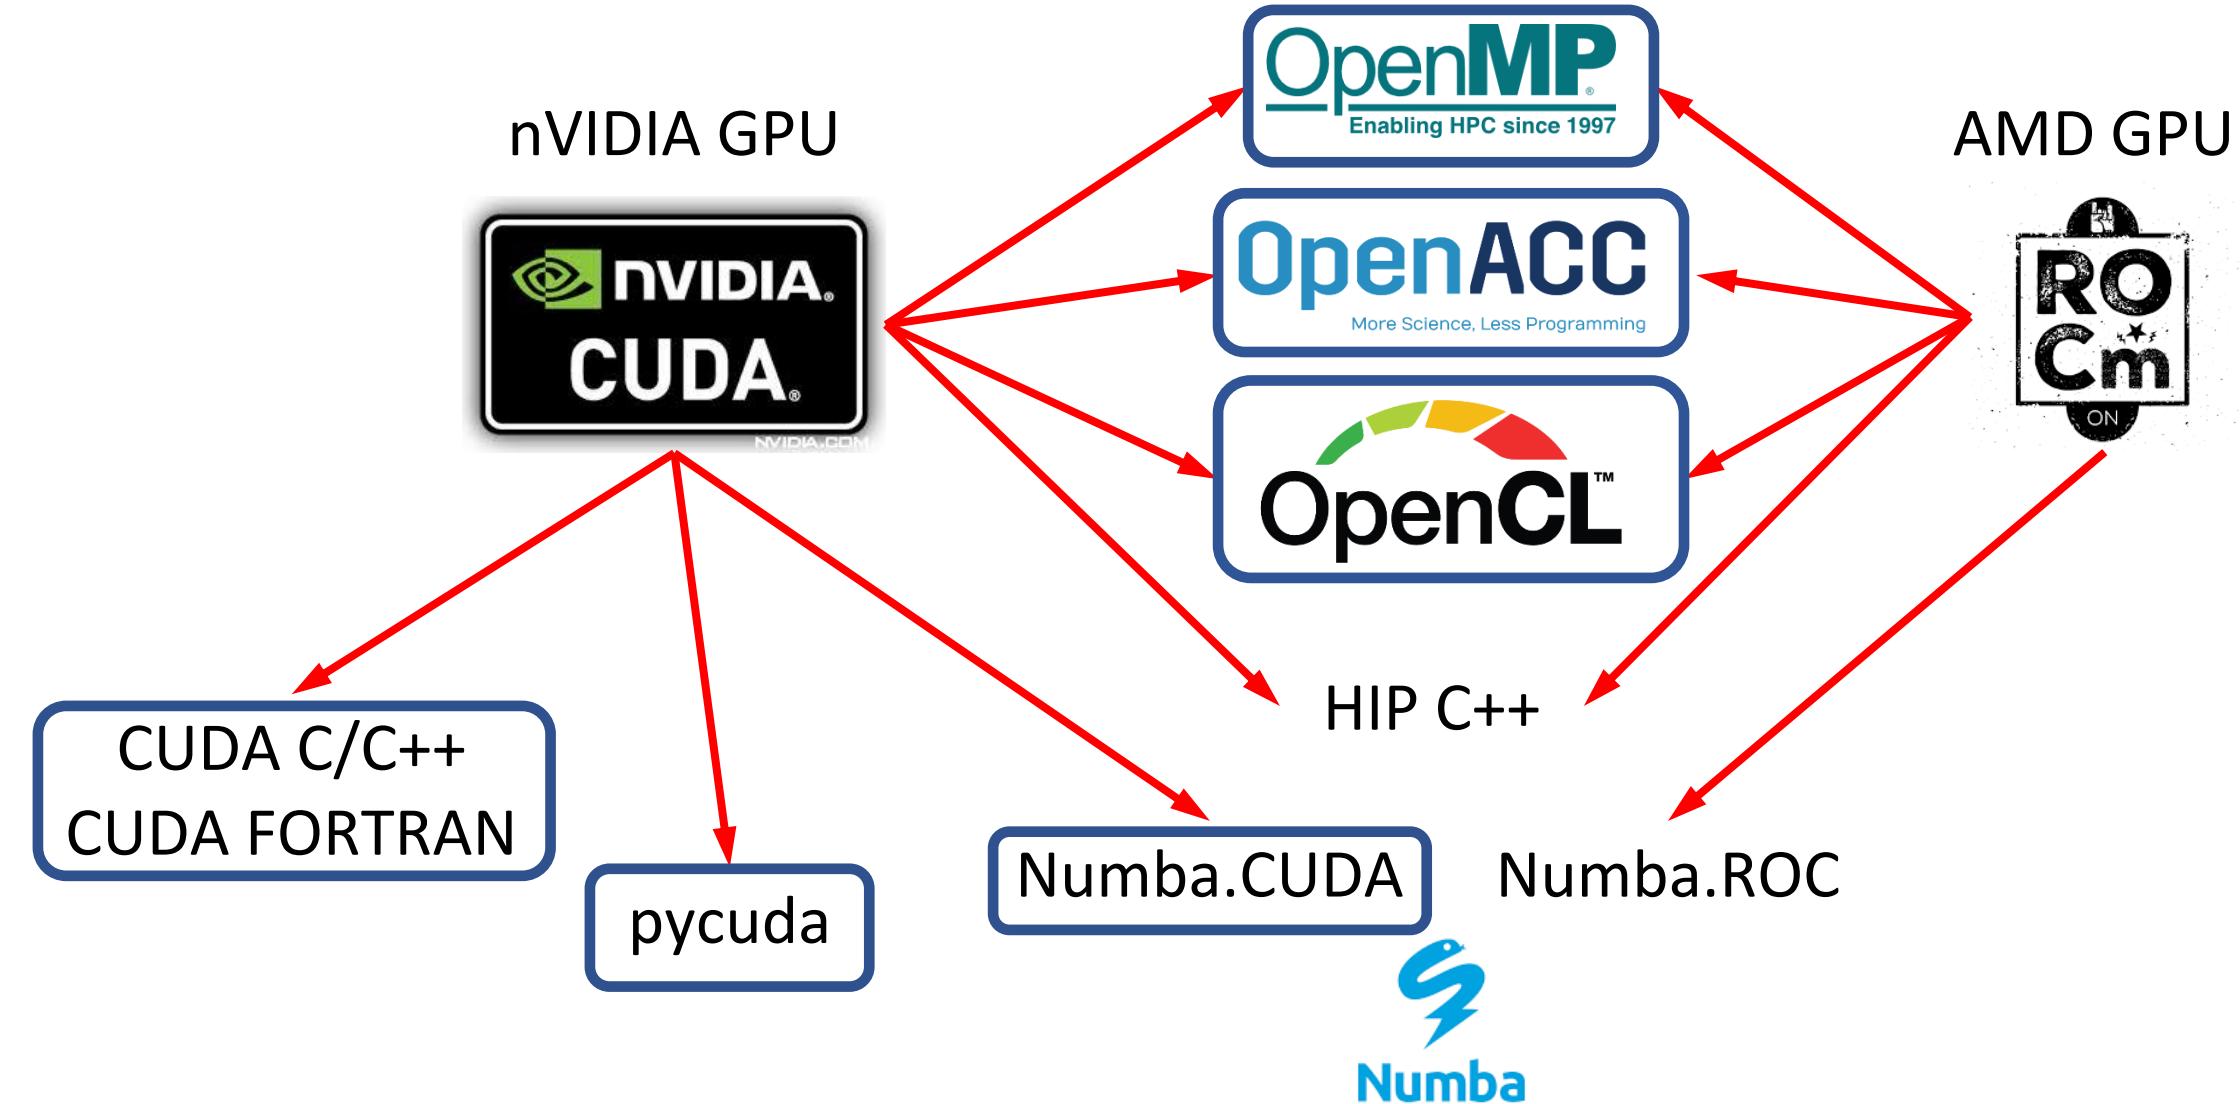


Asynchronous execution helps
to hide the data transfer costs

Writing your kernel : GPU programming paradigms



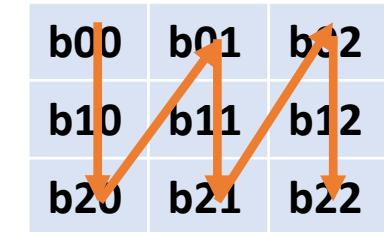
Writing your kernel : GPU programming paradigms



Sub-optimal memory access will kill the performance...

The less efficient way (column-major):

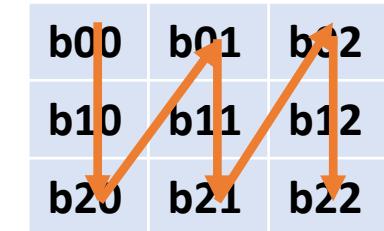
```
for(int i=0; i<n; ++i)
    for(int j=0; j<n; ++j)
        for(int k=0; k<n; ++k)
            c[i*n+j] += a[i*n+k] * b[k*n+j];
```



Sub-optimal memory access will kill the performance...

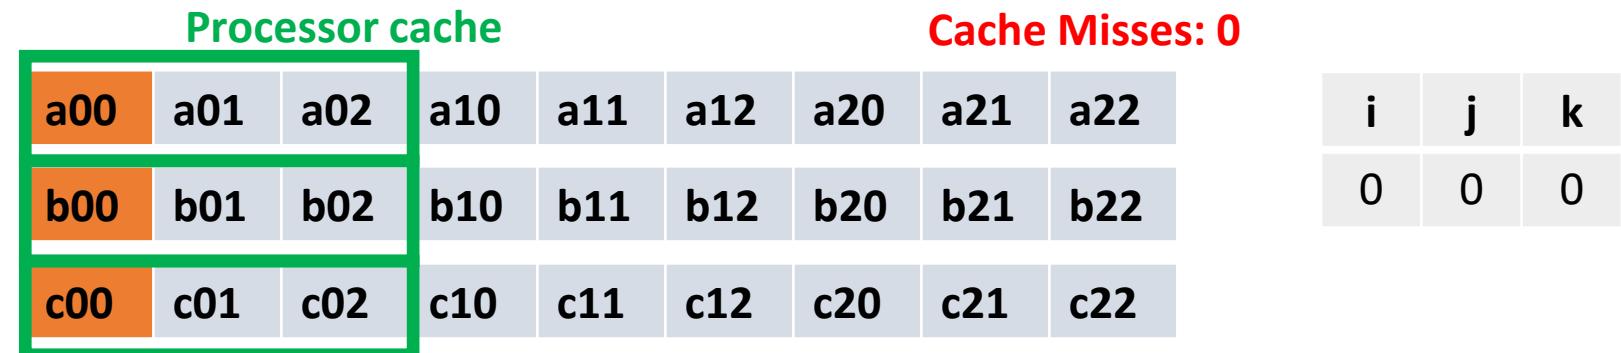
The less efficient way (column-major):

```
for(int i=0; i<n; ++i)
    for(int j=0; j<n; ++j)
        for(int k=0; k<n; ++k)
            c[i*n+j] += a[i*n+k] * b[k*n+j];
```



- Processors **load data into their cache** for fast reuse
- They always load a bunch of data at the same time
- They do **branch-prediction** and **pre-fetching**!

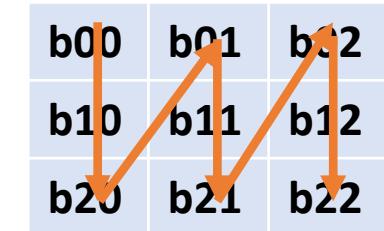
However here we are not helping ...



Sub-optimal memory access will kill the performance...

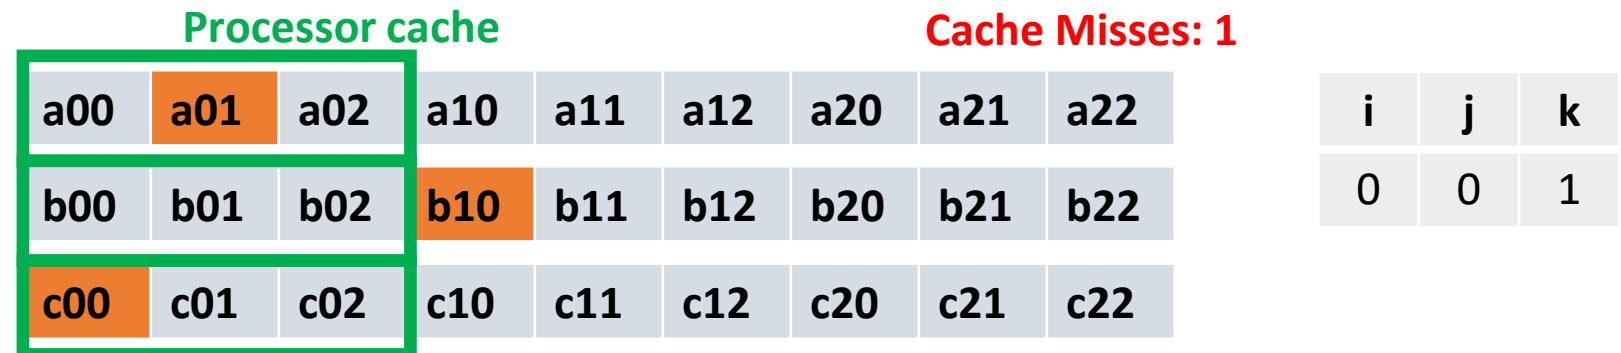
The less efficient way (column-major):

```
for(int i=0; i<n; ++i)
    for(int j=0; j<n; ++j)
        for(int k=0; k<n; ++k)
            c[i*n+j] += a[i*n+k] * b[k*n+j];
```



- Processors **load data into their cache** for fast reuse
- They always load a bunch of data at the same time
- They do **branch-prediction** and **pre-fetching**!

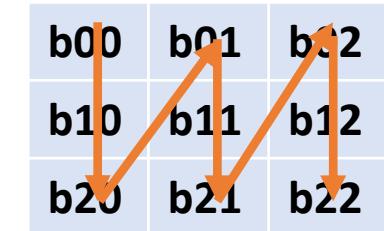
However here we are not helping ...



Sub-optimal memory access will kill the performance...

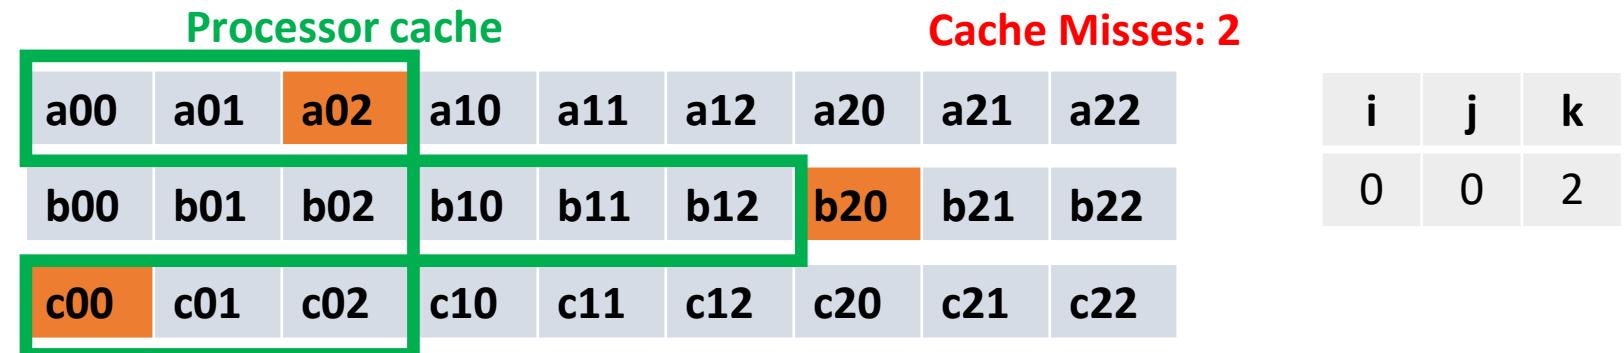
The less efficient way (column-major):

```
for(int i=0; i<n; ++i)
    for(int j=0; j<n; ++j)
        for(int k=0; k<n; ++k)
            c[i*n+j] += a[i*n+k] * b[k*n+j];
```



- Processors **load data into their cache** for fast reuse
- They always load a bunch of data at the same time
- They do **branch-prediction** and **pre-fetching**!

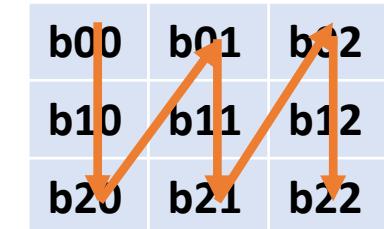
However here we are not helping ...



Sub-optimal memory access will kill the performance...

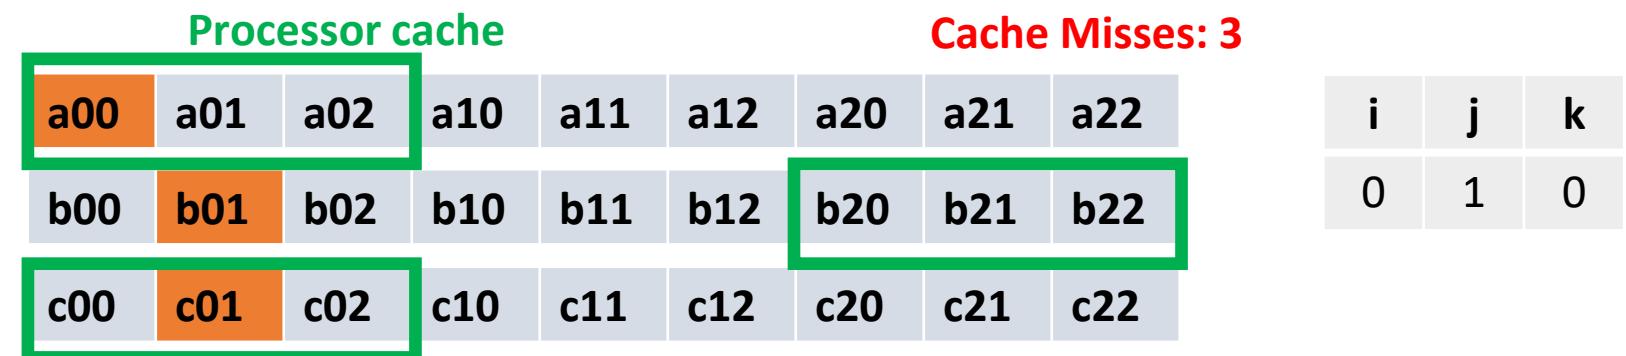
The less efficient way (column-major):

```
for(int i=0; i<n; ++i)
    for(int j=0; j<n; ++j)
        for(int k=0; k<n; ++k)
            c[i*n+j] += a[i*n+k] * b[k*n+j];
```



- Processors **load data into their cache** for fast reuse
- They always load a bunch of data at the same time
- They do **branch-prediction** and **pre-fetching**!

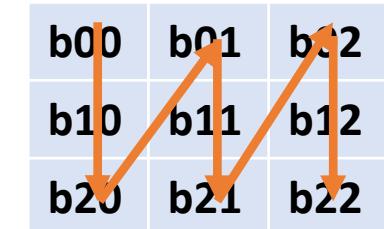
However here we are not helping ...



Sub-optimal memory access will kill the performance...

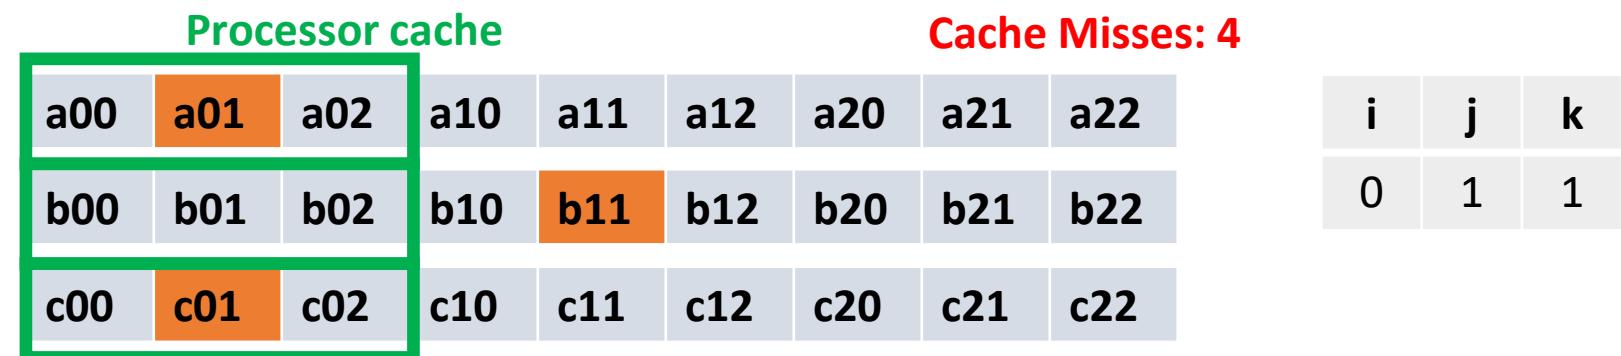
The less efficient way (column-major):

```
for(int i=0; i<n; ++i)
    for(int j=0; j<n; ++j)
        for(int k=0; k<n; ++k)
            c[i*n+j] += a[i*n+k] * b[k*n+j];
```



- Processors **load data into their cache** for fast reuse
- They always load a bunch of data at the same time
- They do **branch-prediction** and **pre-fetching**!

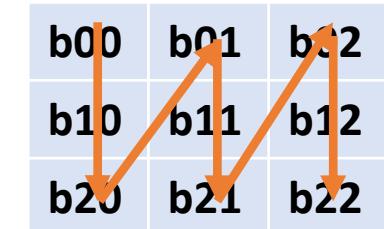
However here we are not helping ...



Sub-optimal memory access will kill the performance...

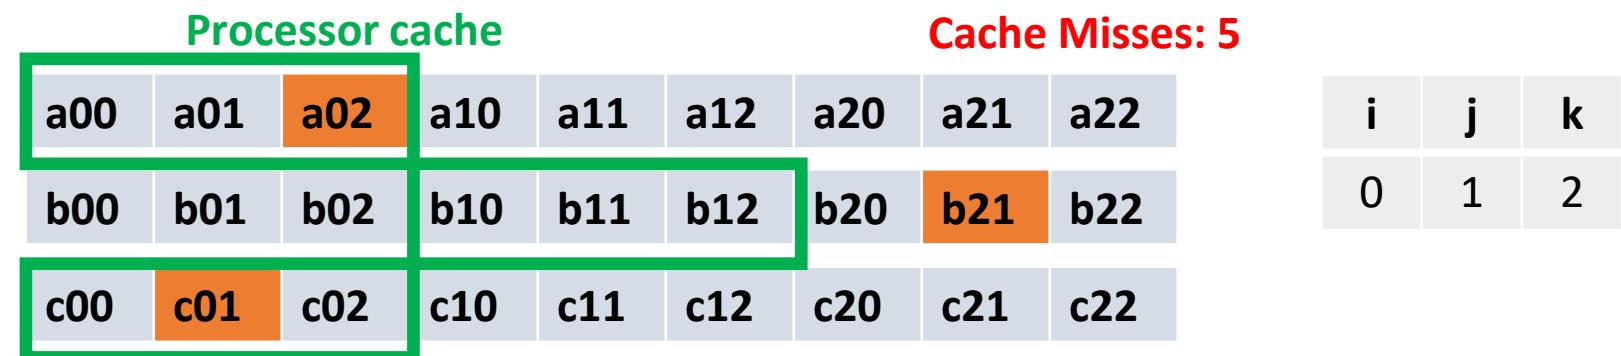
The less efficient way (column-major):

```
for(int i=0; i<n; ++i)
    for(int j=0; j<n; ++j)
        for(int k=0; k<n; ++k)
            c[i*n+j] += a[i*n+k] * b[k*n+j];
```



- Processors **load data into their cache** for fast reuse
- They always load a bunch of data at the same time
- They do **branch-prediction** and **pre-fetching**!

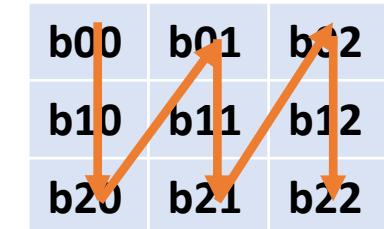
However here we are not helping ...



Sub-optimal memory access will kill the performance...

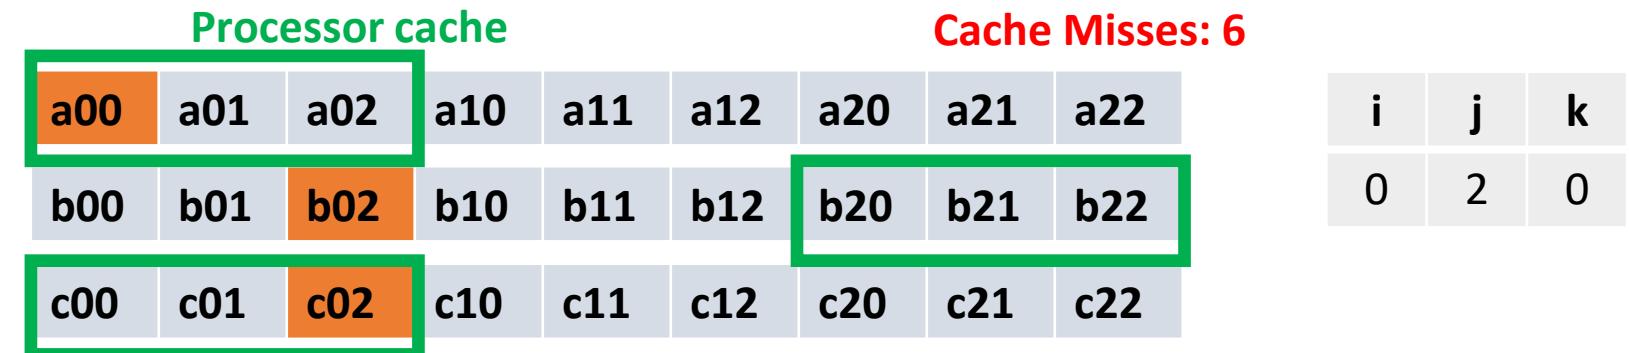
The less efficient way (column-major):

```
for(int i=0; i<n; ++i)
    for(int j=0; j<n; ++j)
        for(int k=0; k<n; ++k)
            c[i*n+j] += a[i*n+k] * b[k*n+j];
```



- Processors **load data into their cache** for fast reuse
- They always load a bunch of data at the same time
- They do **branch-prediction** and **pre-fetching**!

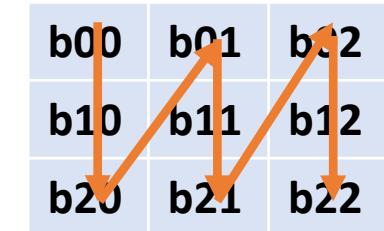
However here we are not helping ...



Sub-optimal memory access will kill the performance...

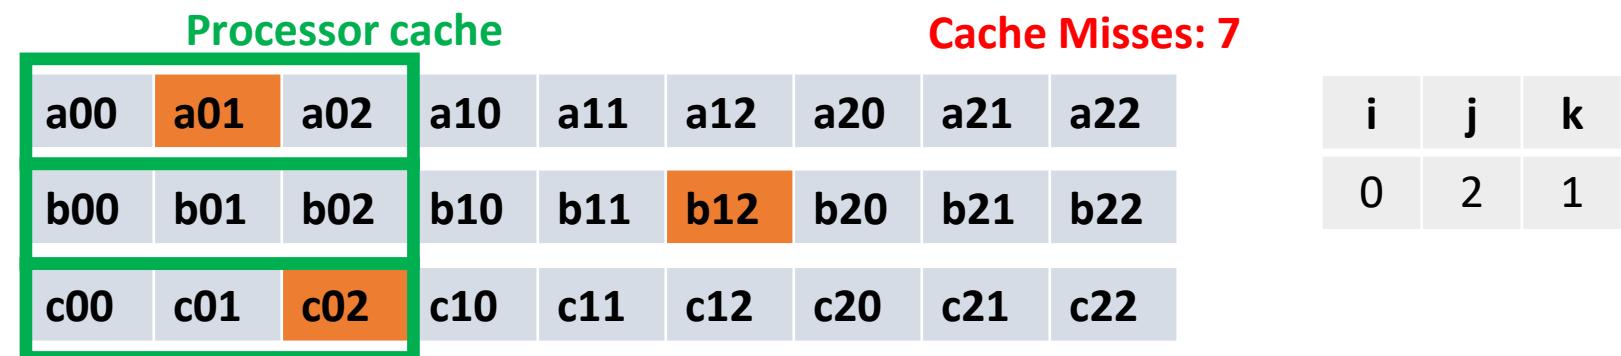
The less efficient way (column-major):

```
for(int i=0; i<n; ++i)
    for(int j=0; j<n; ++j)
        for(int k=0; k<n; ++k)
            c[i*n+j] += a[i*n+k] * b[k*n+j];
```



- Processors **load data into their cache** for fast reuse
- They always load a bunch of data at the same time
- They do **branch-prediction** and **pre-fetching**!

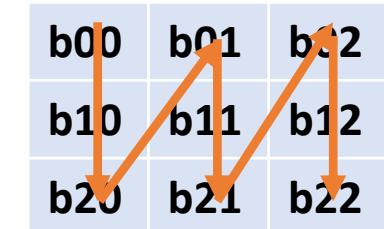
However here we are not helping ...



Sub-optimal memory access will kill the performance...

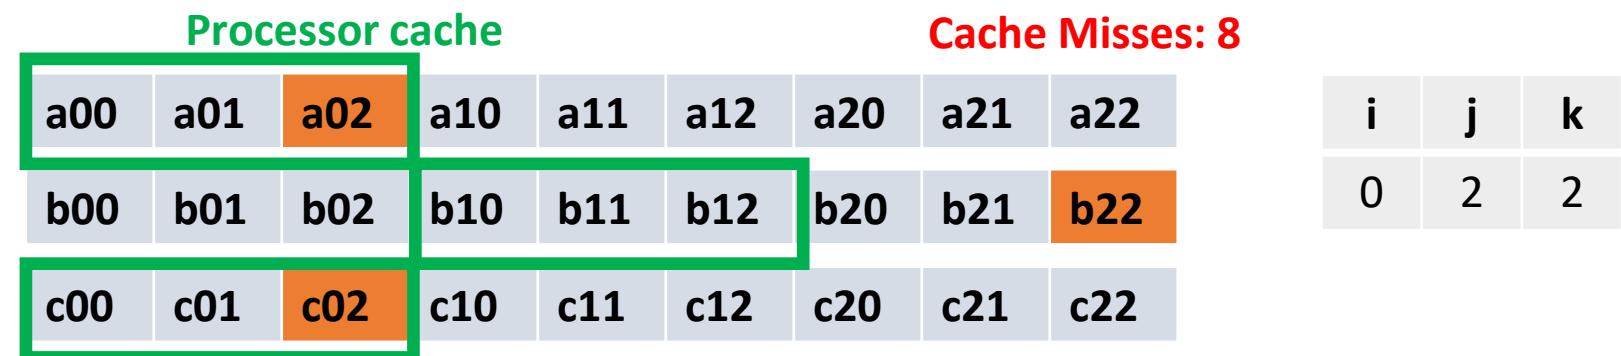
The less efficient way (column-major):

```
for(int i=0; i<n; ++i)
    for(int j=0; j<n; ++j)
        for(int k=0; k<n; ++k)
            c[i*n+j] += a[i*n+k] * b[k*n+j];
```



- Processors **load data into their cache** for fast reuse
- They always load a bunch of data at the same time
- They do **branch-prediction** and **pre-fetching**!

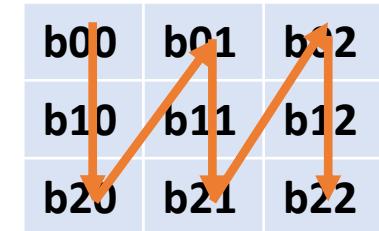
However here we are not helping ...



Sub-optimal memory access will kill the performance...

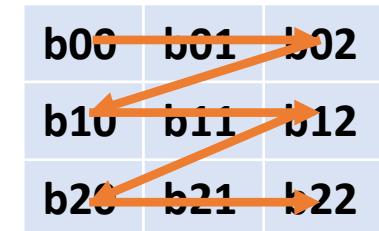
The less efficient way (column-major):

```
for(int i=0; i<n; ++i)
    for(int j=0; j<n; ++j)
        for(int k=0; k<n; ++k)
            c[i*n+j] += a[i*n+k] * b[k*n+j];
```



The better way (row-major):

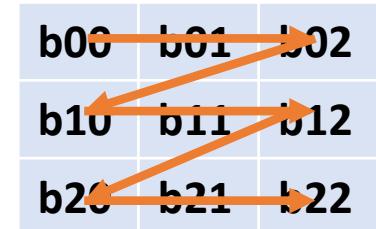
```
for(int i=0; i<n; ++i)
    for(int k=0; k<n; ++k)
        for(int j=0; j<n; ++j)
            c[i*n+j] += a[i*n+k] * b[k*n+j];
```



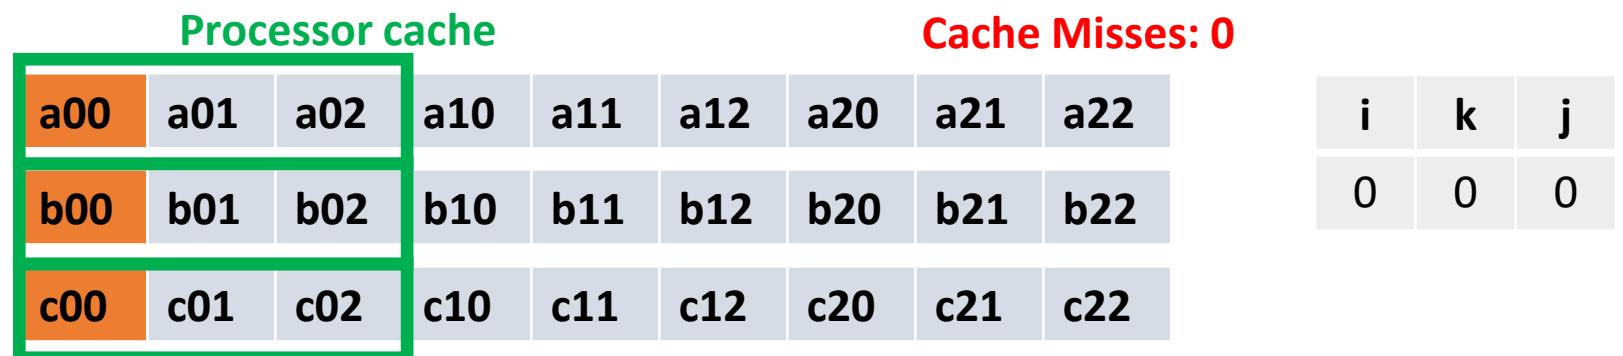
**Here, the most inner loop can
be executed in parallel** (without
write conflicts on matrix c)

The better way (row-major):

```
for(int i=0; i<n; ++i)
    for(int k=0; k<n; ++k)
        for(int j=0; j<n; ++j)
            c[i*n+j] += a[i*n+k] * b[k*n+j];
```

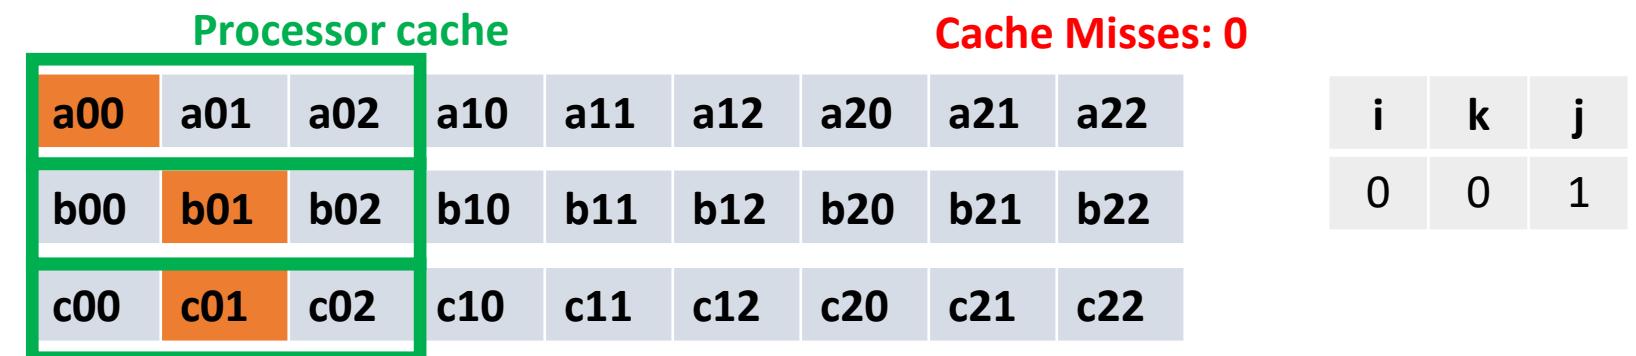
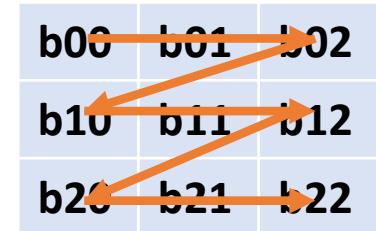


Let's go...



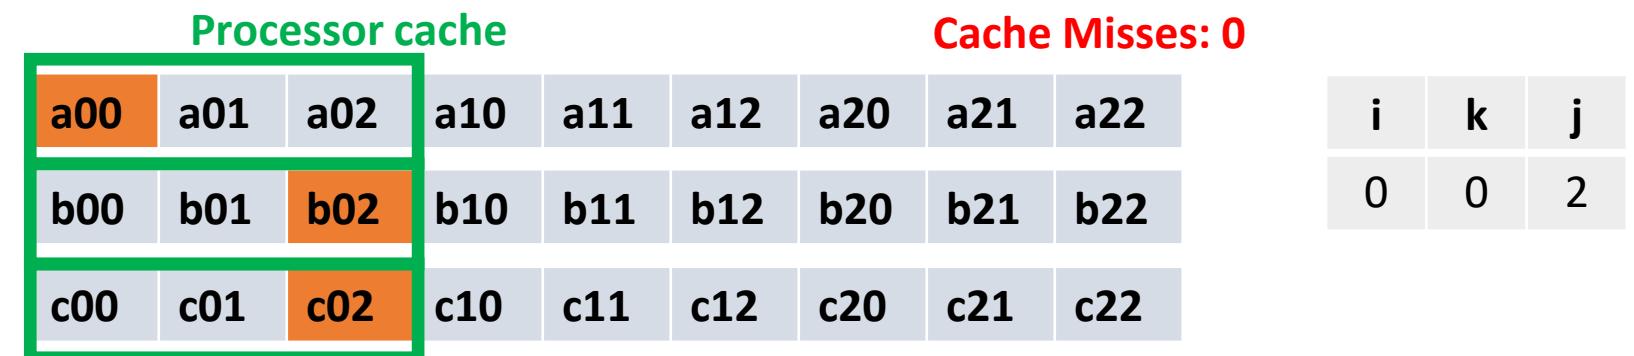
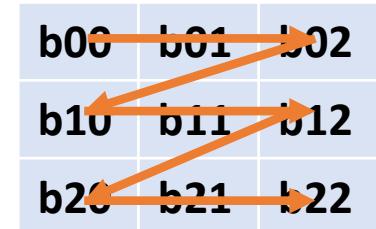
The better way (row-major):

```
for(int i=0; i<n; ++i)
    for(int k=0; k<n; ++k)
        for(int j=0; j<n; ++j)
            c[i*n+j] += a[i*n+k] * b[k*n+j];
```



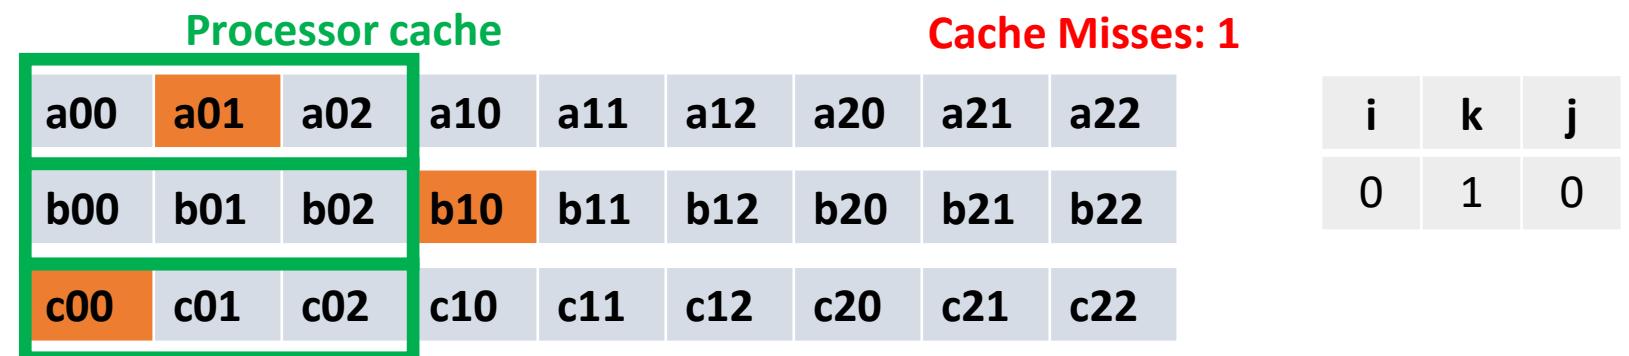
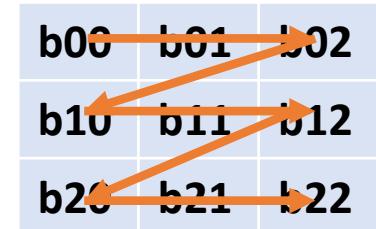
The better way (row-major):

```
for(int i=0; i<n; ++i)
    for(int k=0; k<n; ++k)
        for(int j=0; j<n; ++j)
            c[i*n+j] += a[i*n+k] * b[k*n+j];
```



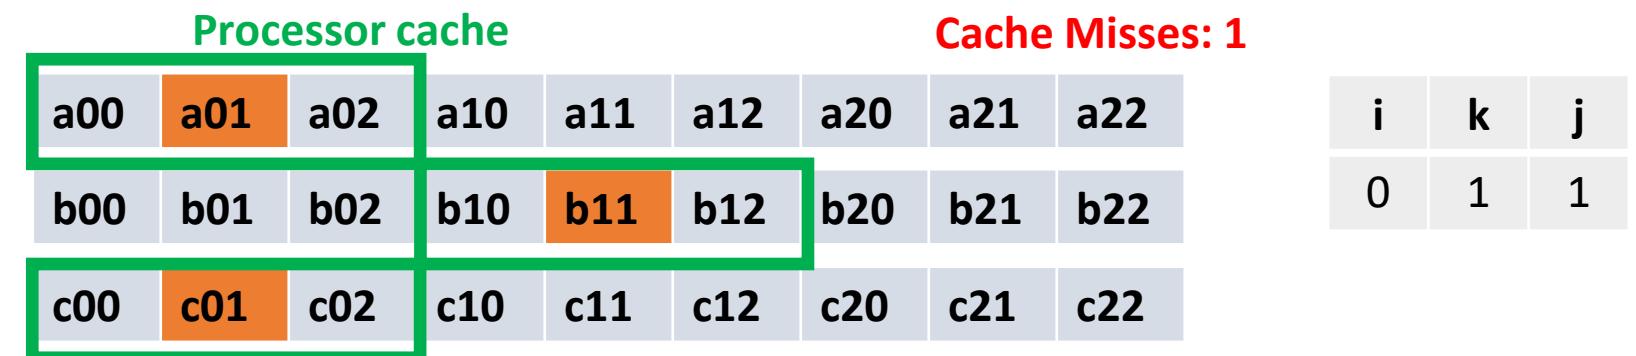
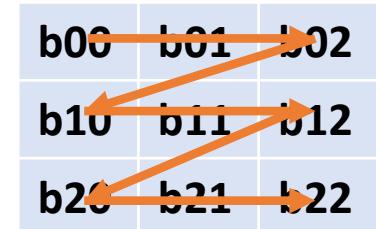
The better way (row-major):

```
for(int i=0; i<n; ++i)
    for(int k=0; k<n; ++k)
        for(int j=0; j<n; ++j)
            c[i*n+j] += a[i*n+k] * b[k*n+j];
```



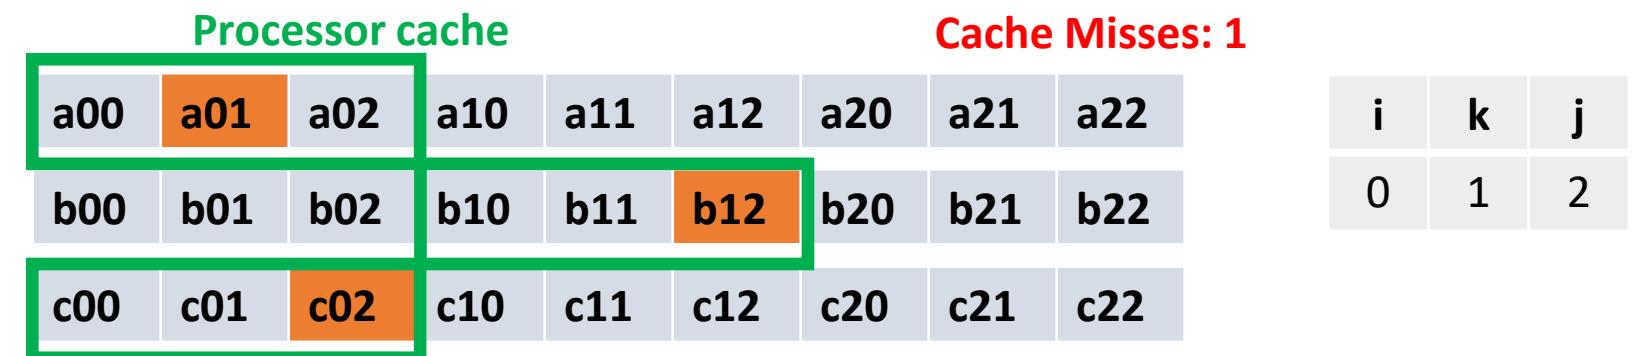
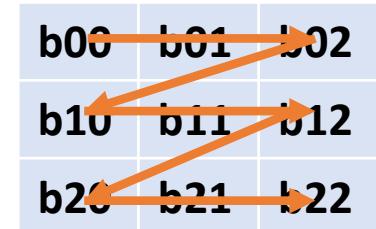
The better way (row-major):

```
for(int i=0; i<n; ++i)
    for(int k=0; k<n; ++k)
        for(int j=0; j<n; ++j)
            c[i*n+j] += a[i*n+k] * b[k*n+j];
```



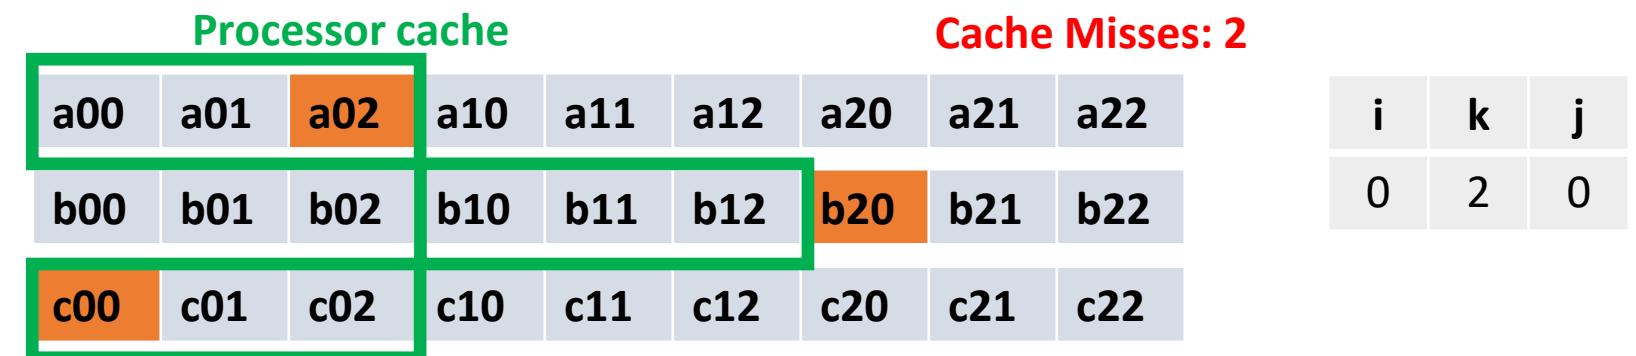
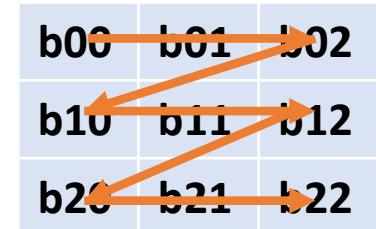
The better way (row-major):

```
for(int i=0; i<n; ++i)
    for(int k=0; k<n; ++k)
        for(int j=0; j<n; ++j)
            c[i*n+j] += a[i*n+k] * b[k*n+j];
```



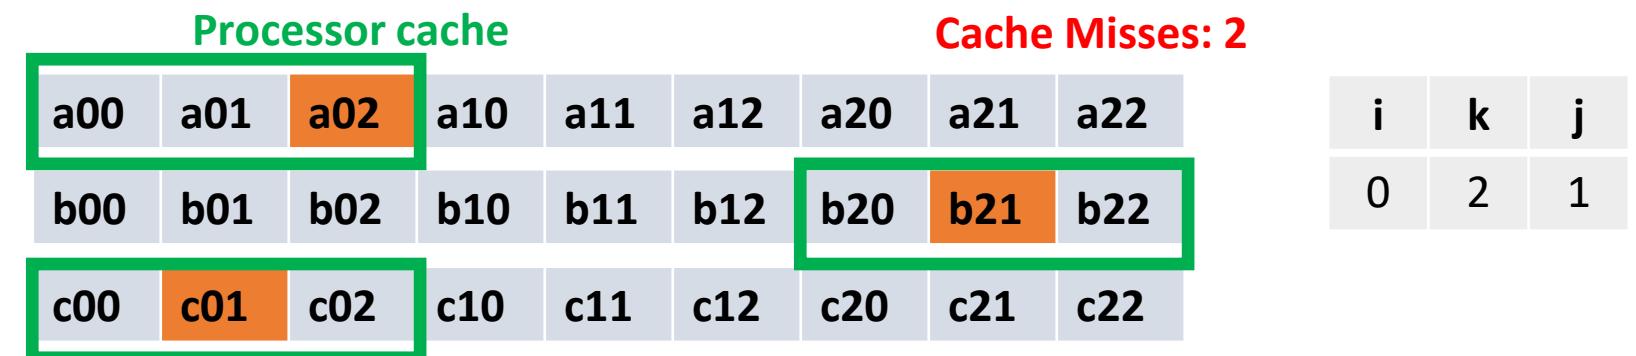
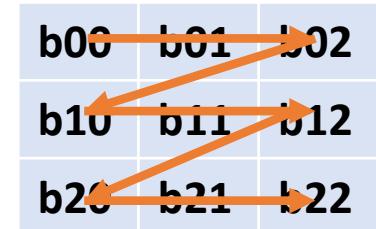
The better way (row-major):

```
for(int i=0; i<n; ++i)
    for(int k=0; k<n; ++k)
        for(int j=0; j<n; ++j)
            c[i*n+j] += a[i*n+k] * b[k*n+j];
```



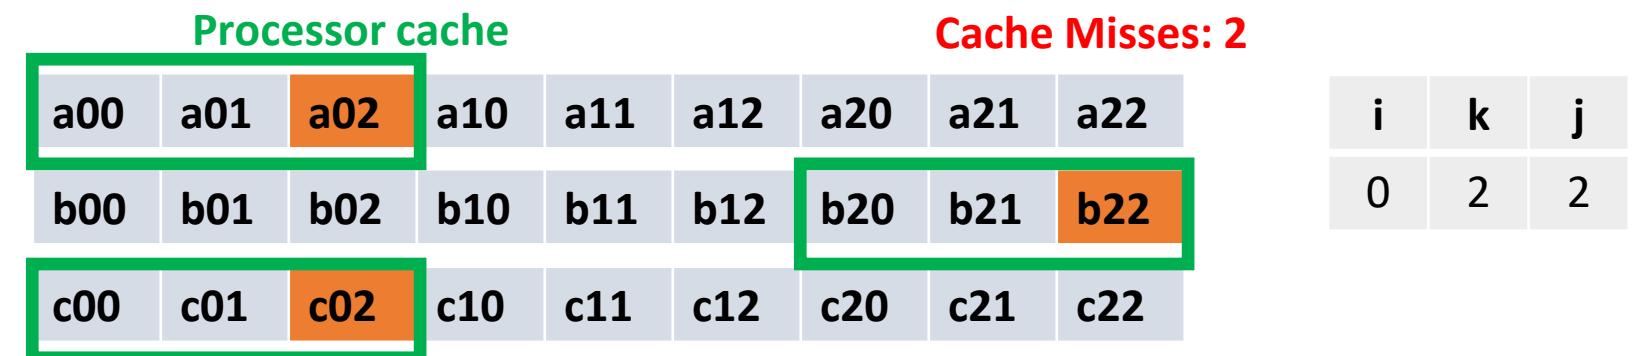
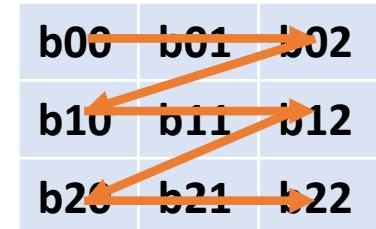
The better way (row-major):

```
for(int i=0; i<n; ++i)
    for(int k=0; k<n; ++k)
        for(int j=0; j<n; ++j)
            c[i*n+j] += a[i*n+k] * b[k*n+j];
```



The better way (row-major):

```
for(int i=0; i<n; ++i)
    for(int k=0; k<n; ++k)
        for(int j=0; j<n; ++j)
            c[i*n+j] += a[i*n+k] * b[k*n+j];
```



Matrix multiplication example

- We create as many kernel instances as we need!
- In this case, we divide the problems into 2D blocks of 16*16 threads each
- The grid of blocks is large enough to accommodate all the blocks

```
__global__ void cuda_mul(float* a, float* b, float* c, int size) {
    int row = blockIdx.y*blockDim.y+threadIdx.y;
    int col = blockIdx.x*blockDim.x+threadIdx.x;
    for (int i = 0; i < size; i++)
        c[row*size+col] += a[row*size+i] * b[i*size+col];
}
```

CUDA Kernel

```
cudaMemcpy(dm1, a, sizeof(float)*size*size, cudaMemcpyHostToDevice);
cudaMemcpy(dm2, b, sizeof(float)*size*size, cudaMemcpyHostToDevice);
cudaMemcpy(dm3, c, sizeof(float)*size*size, cudaMemcpyHostToDevice);

dim3 blockSize = dim3(16, 16);
dim3 gridSize = dim3(size / blockSize.x, size/ blockSize.y);

cuda_mul<<<gridSize, blockSize>>>(dm1, dm2, dm3, size);

cudaMemcpy(c, dm3, sizeof(float)*size*size, cudaMemcpyDeviceToHost);
```

Kernel Call

Matrix multiplication example

Each thread inside a 2D block receives:

- blockIdx.x, blockIdx.y
- threadIdx.x, threadIdx.y

This information is used to retrieve the row and col indices (formerly i and j).

This is the only information that differs between threads!

Remember, the data is not sent to individual threads

Instead, the data is mapped onto the grid of blocks

And threads know which data to access thanks to the block and thread ids.

```
__global__ void cuda_mul(float* a, float* b, float* c, int size) {  
    int row = blockIdx.y*blockDim.y+threadIdx.y;  
    int col = blockIdx.x*blockDim.x+threadIdx.x;  
    for (int i = 0; i < size; i++)  
        c[row*size+col] += a[row*size+i] * b[i*size+col];  
}
```

CUDA Kernel



Matrix multiplication example

Each thread inside a 2D block receives:

- blockIdx.x, blockIdx.y
- threadIdx.x, threadIdx.y

This information is used to retrieve the row and col indices (formerly i and j).

This is the only information that differs between threads!

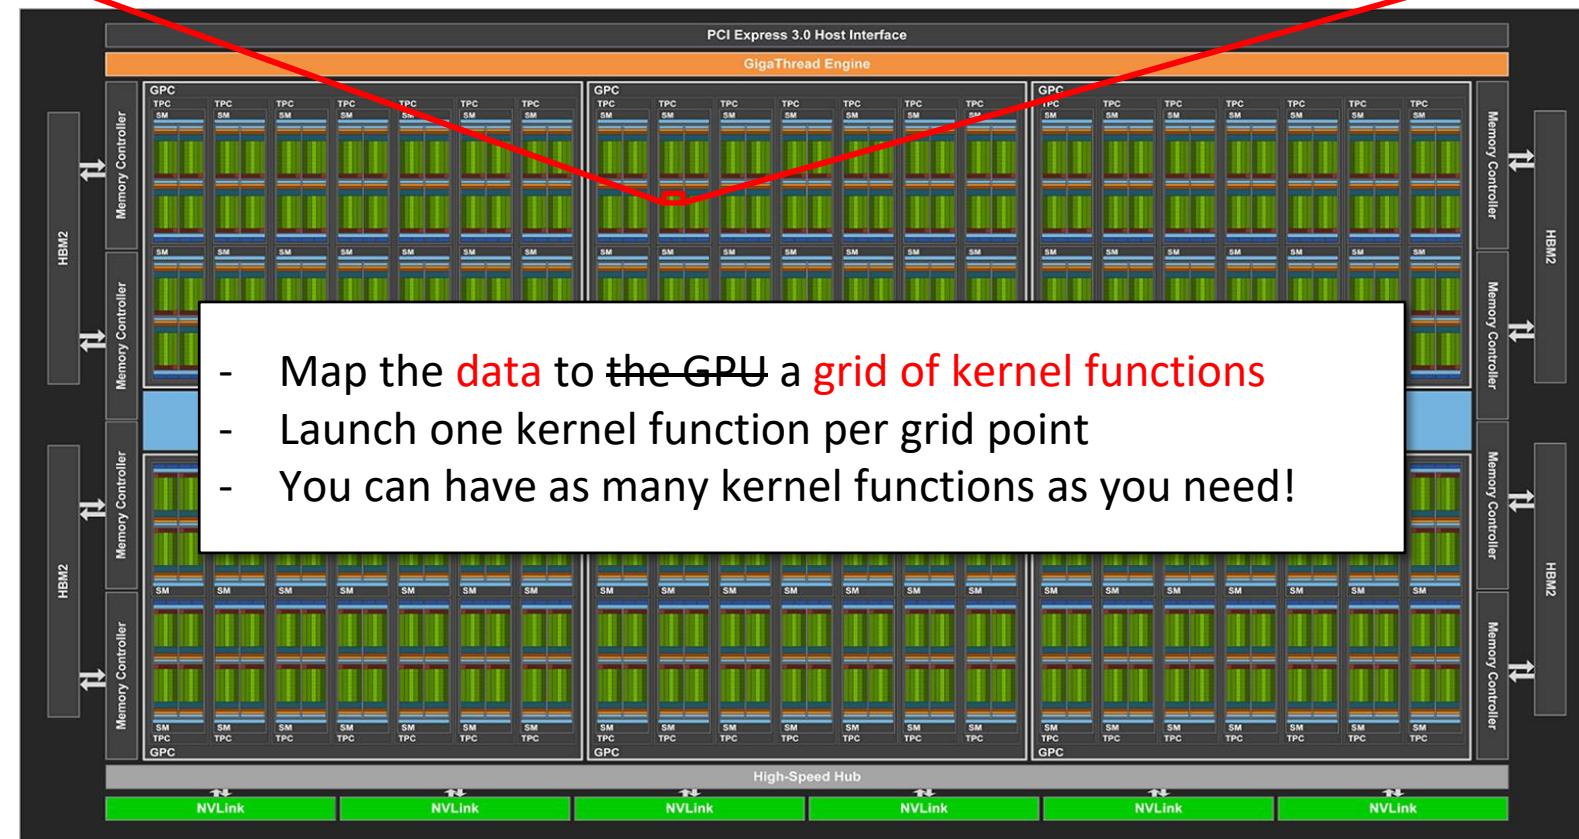
Remember, the data is not sent to individual threads

Instead, the data is mapped onto the grid of blocks

And threads know which data to access thanks to the block and thread ids.

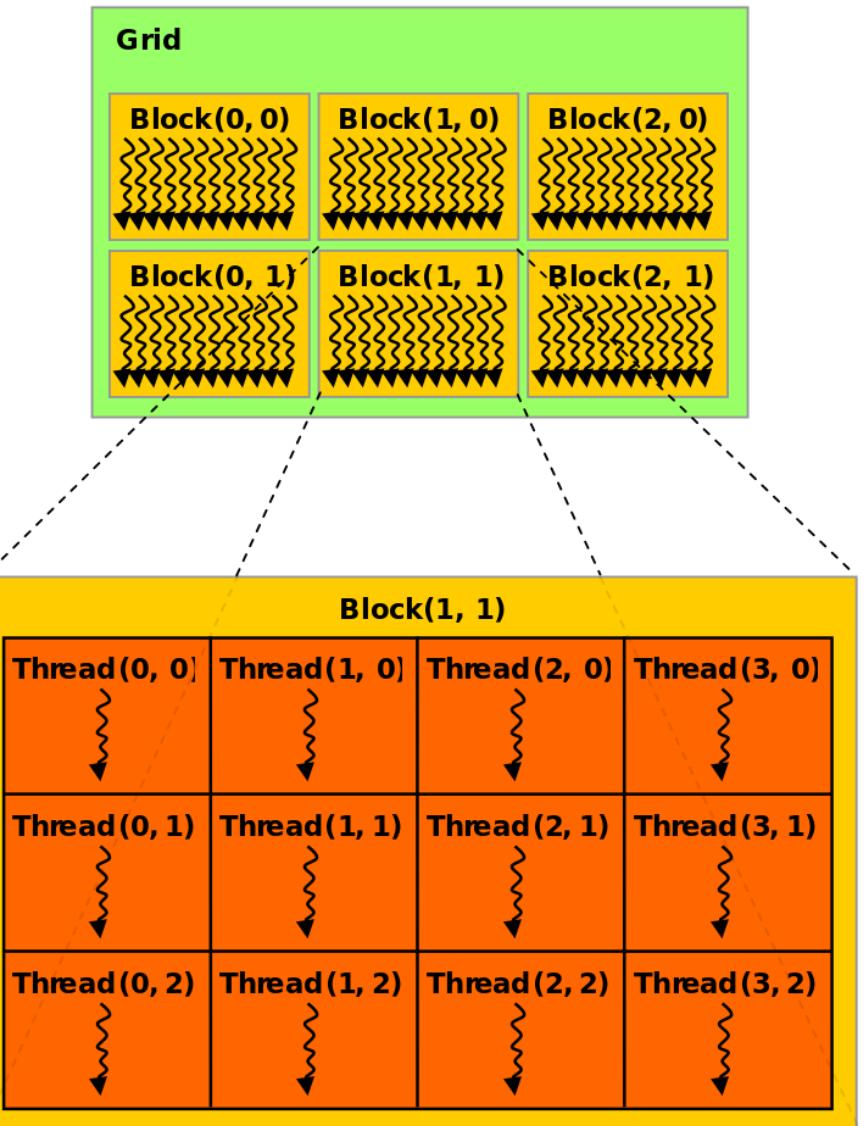
```
__global__ void cuda_mul(float* a, float* b, float* c, int size) {  
    int row = blockIdx.y*blockDim.y+threadIdx.y;  
    int col = blockIdx.x*blockDim.x+threadIdx.x;  
    for (int i = 0; i < size; i++)  
        c[row*size+col] += a[row*size+i] * b[i*size+col];  
}
```

CUDA Kernel



Reminder: Grids, blocks, threads and warps

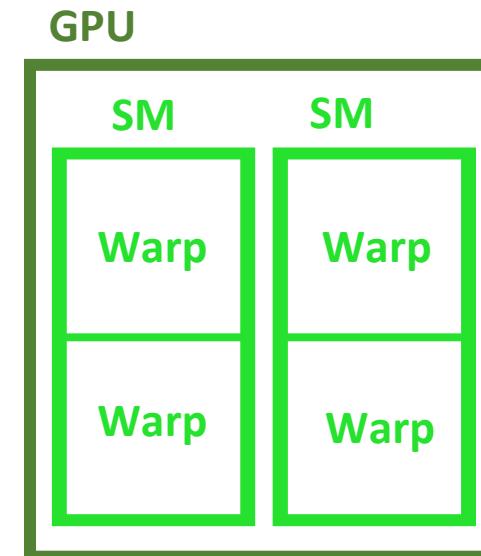
- Grid:
 - 1 to as many blocks as you want
- Block
 - 1 to 1024 threads
 - Should be a multiple of 32 threads
 - Execute on a single Streaming Multiprocessor (SM)
 - Possibility of having shared memory between threads
 - Is executed in **warps** of 32 threads
- Warps
 - Warps always execute 32 threads inside a SM
 - **All the 32 threads in a warp execute the same instruction (lockstep)**



Execution of a block on a GPU

Grid (3x2 blocks)

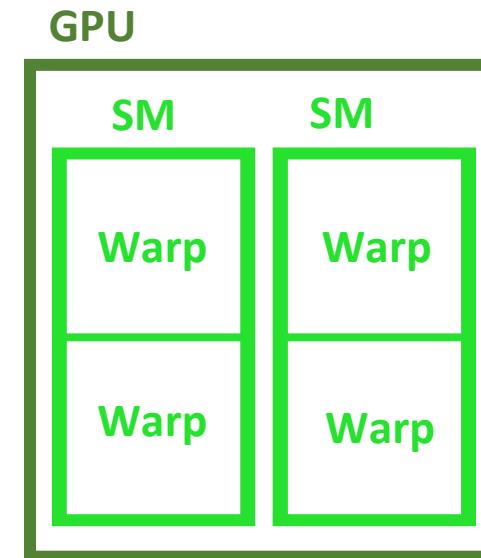
Block (0, 0) 128 threads	Block (0, 1) 128 threads	Block (0, 2) 128 threads
Block (1, 0) 128 threads	Block (1, 1) 128 threads	Block (1, 2) 128 threads



Execution of a block on a GPU

Grid (3x2 blocks)

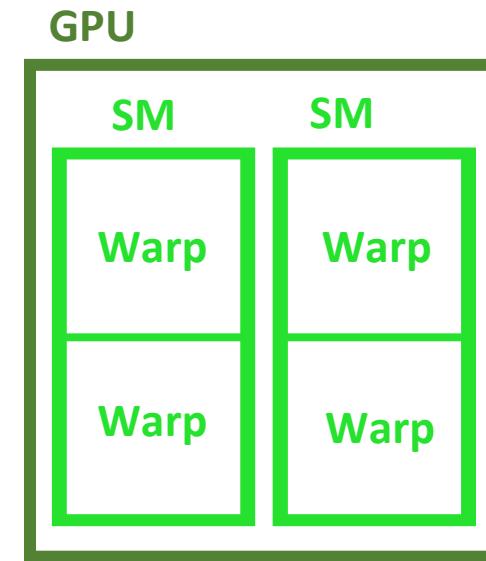
Block (0, 0) 128 threads	Block (0, 1) 128 threads	Block (0, 2) 128 threads
Block (1, 0) 128 threads	Block (1, 1) 128 threads	Block (1, 2) 128 threads



Execution of a block on a GPU

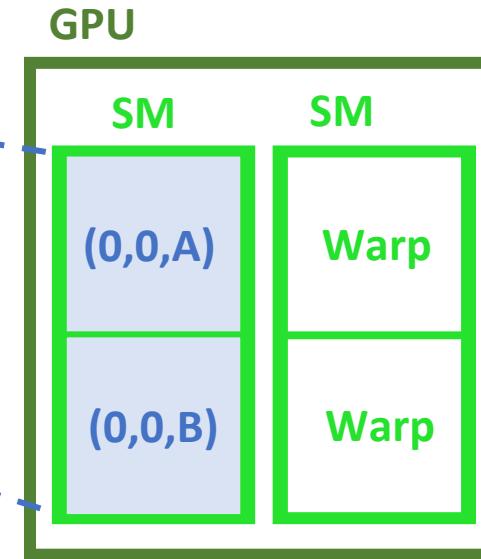
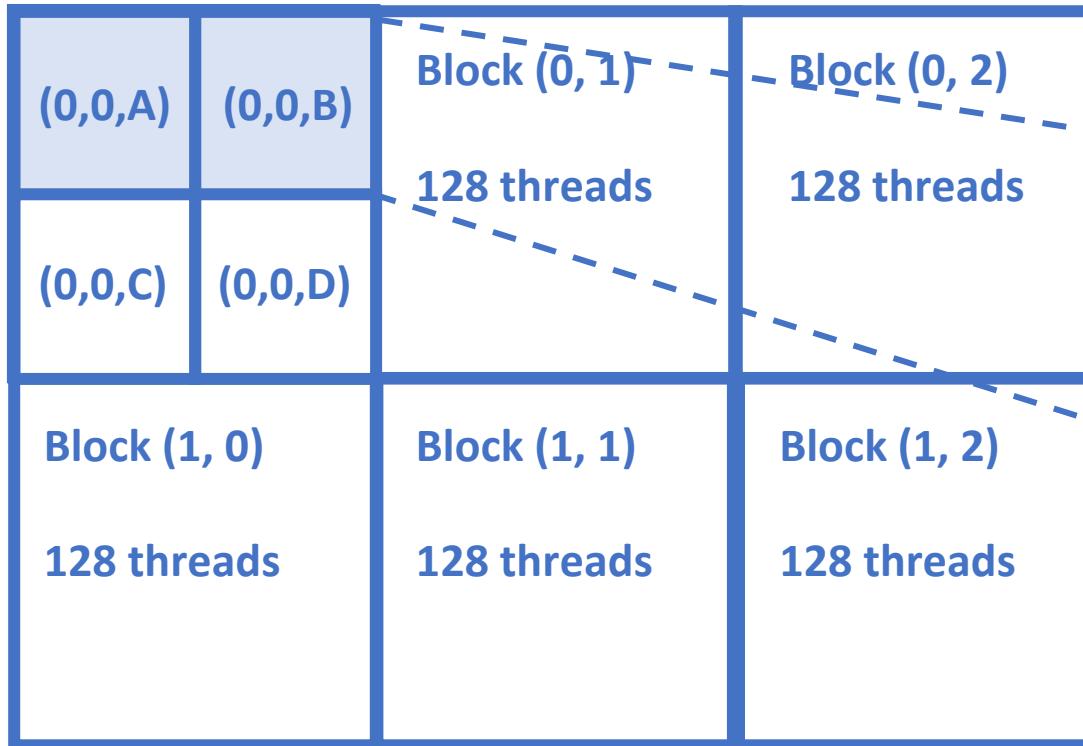
Grid (3x2 blocks)

32 threads	32 threads	Block (0, 1) 128 threads	Block (0, 2) 128 threads
32 threads	32 threads		
Block (1, 0) 128 threads	Block (1, 1) 128 threads	Block (1, 2) 128 threads	



Execution of a block on a GPU

Grid (3x2 blocks)

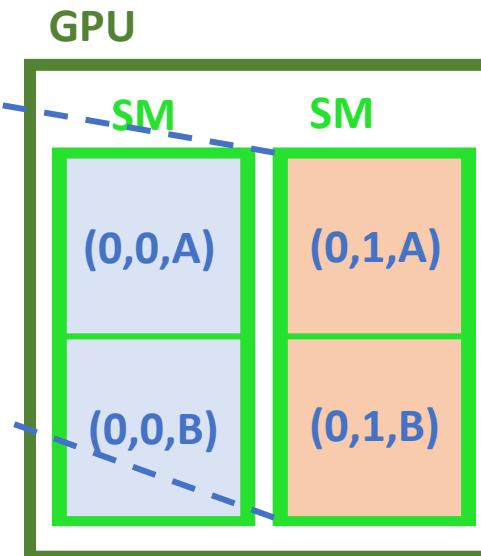
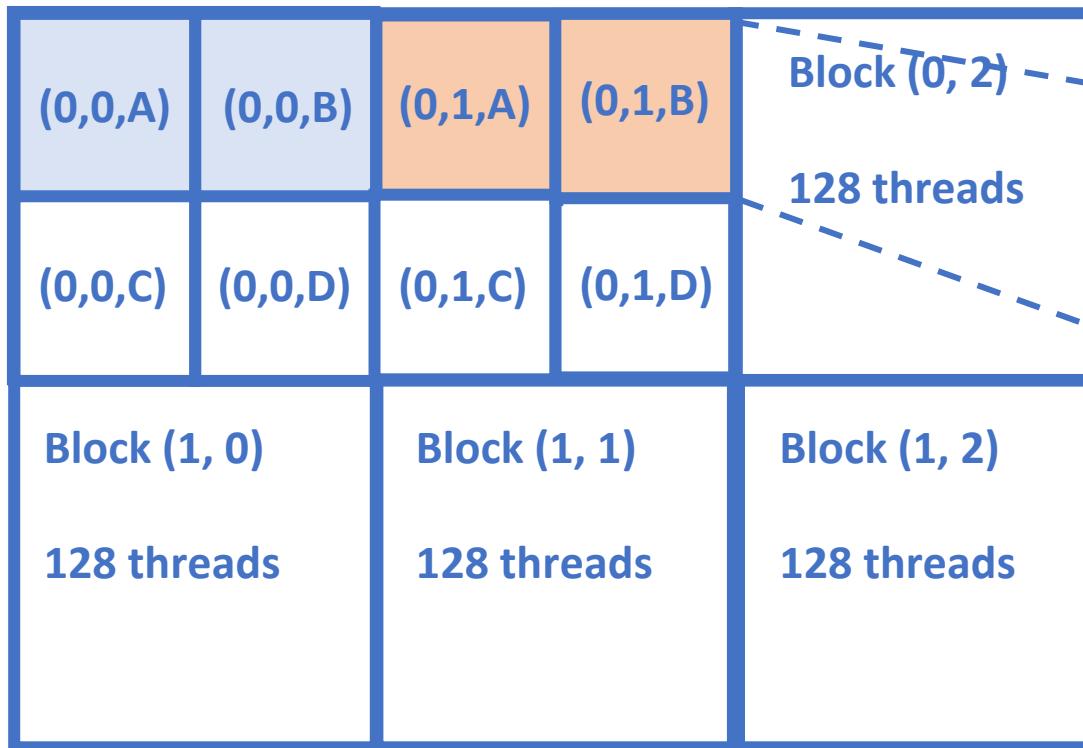


Remember: a block can only be executed by a single SM!

- The first block is scheduled on two warps of 32 threads on the first SM.

Execution of a block on a GPU

Grid (3x2 blocks)

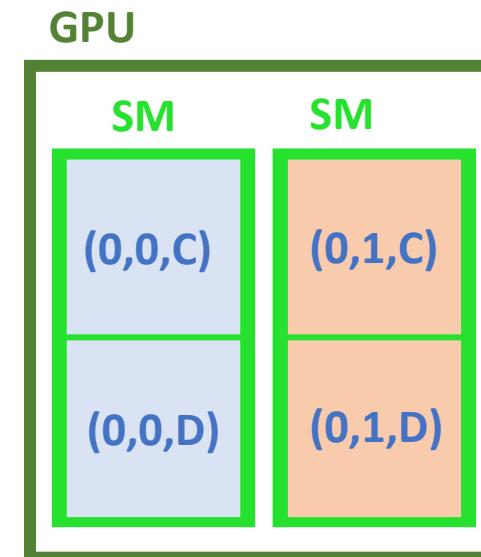
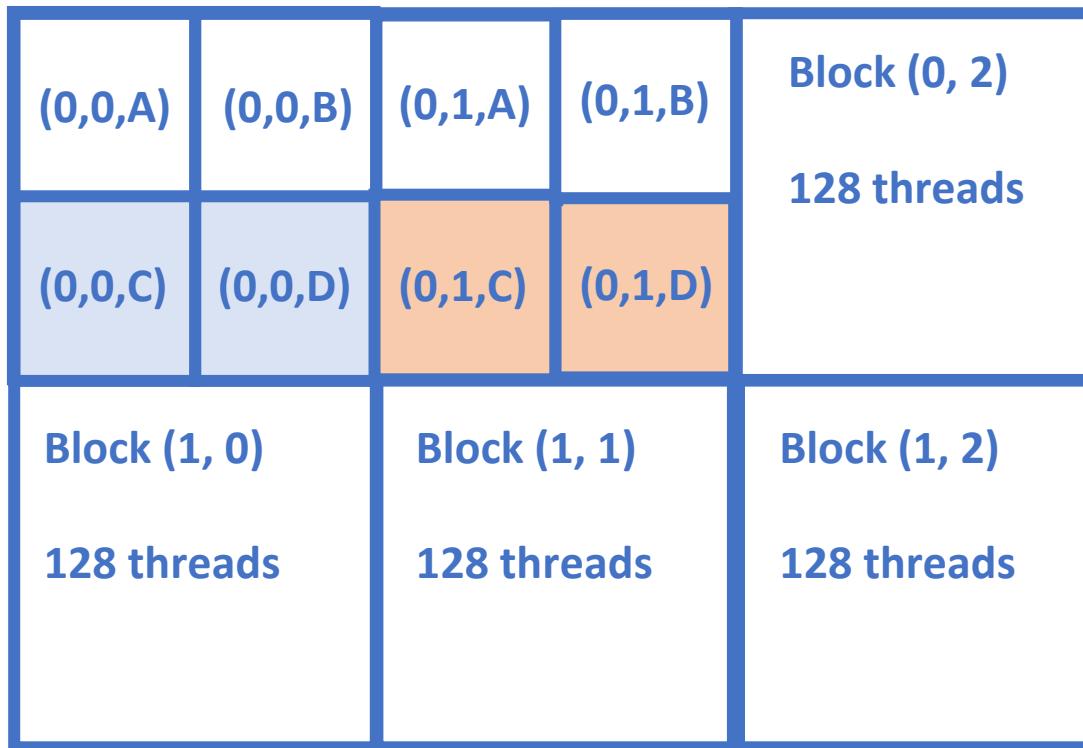


Remember: a block can only be executed by a single SM!

- The first block is scheduled on two warps of 32 threads on the first SM.
- The second block is scheduled on the second warp and the GPU is occupied at 100%

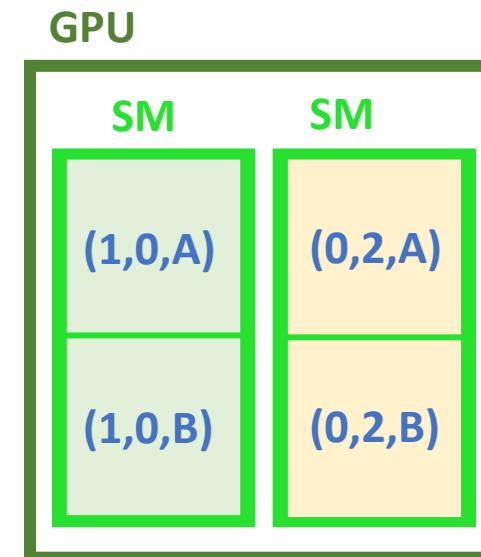
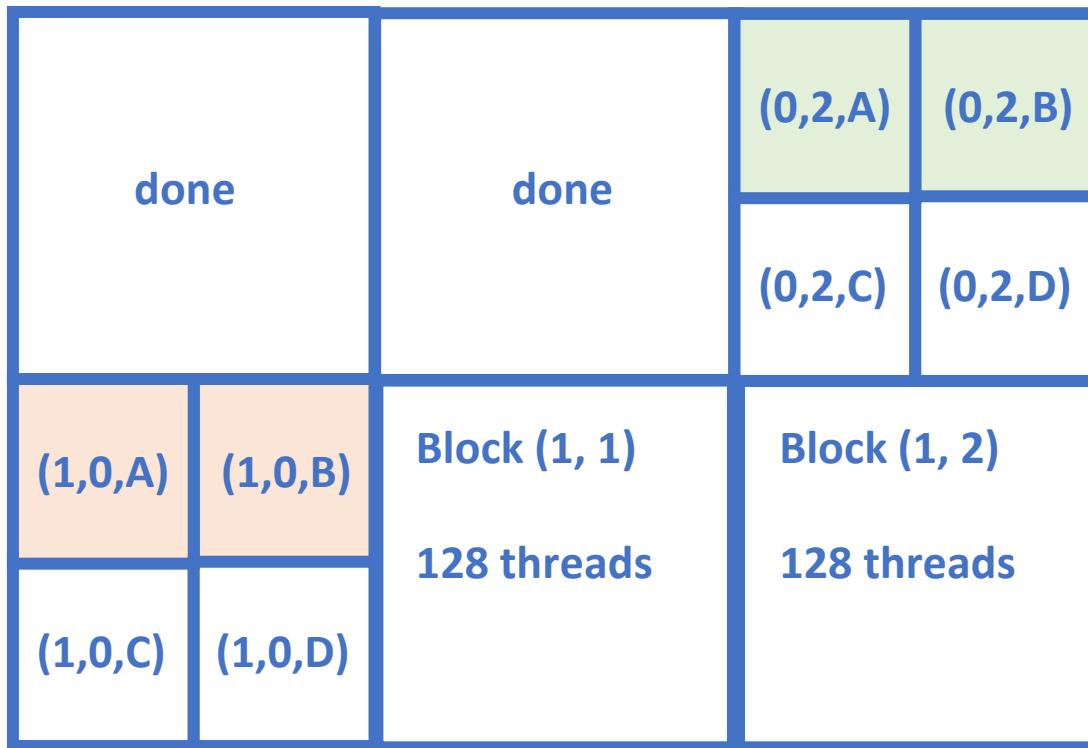
Execution of a block on a GPU

Grid (3x2 blocks)



Execution of a block on a GPU

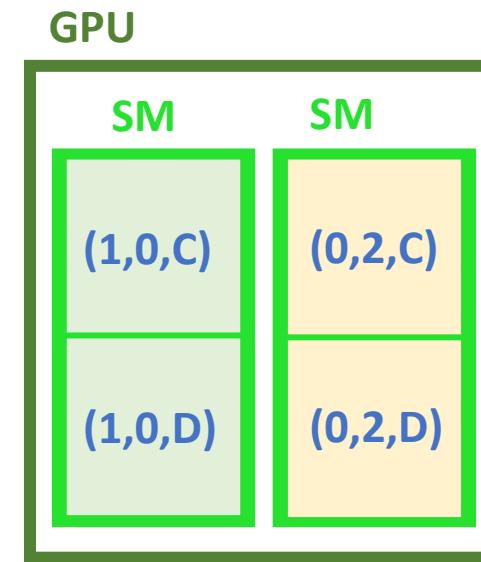
Grid (3x2 blocks)



As soon as blocks (0,0) and (0,1) are done, the GPU will start executing the next blocks

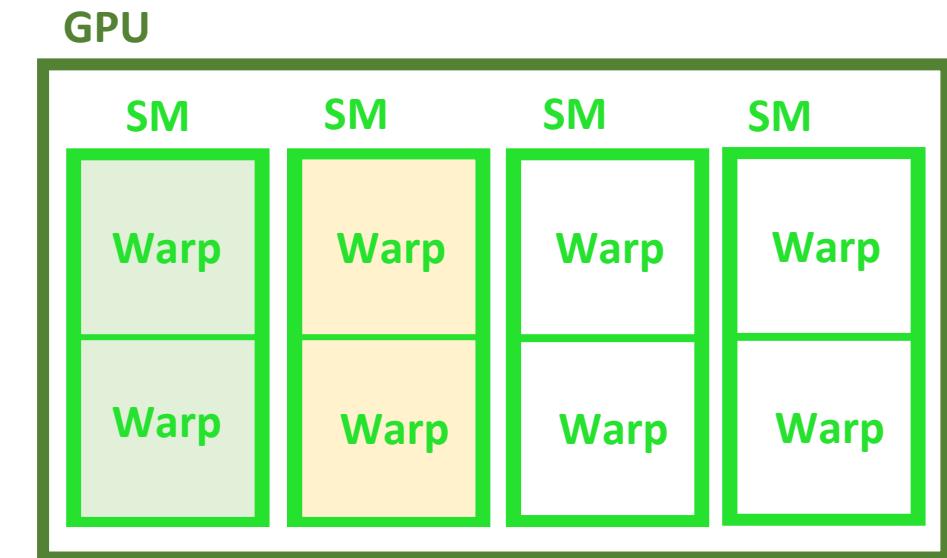
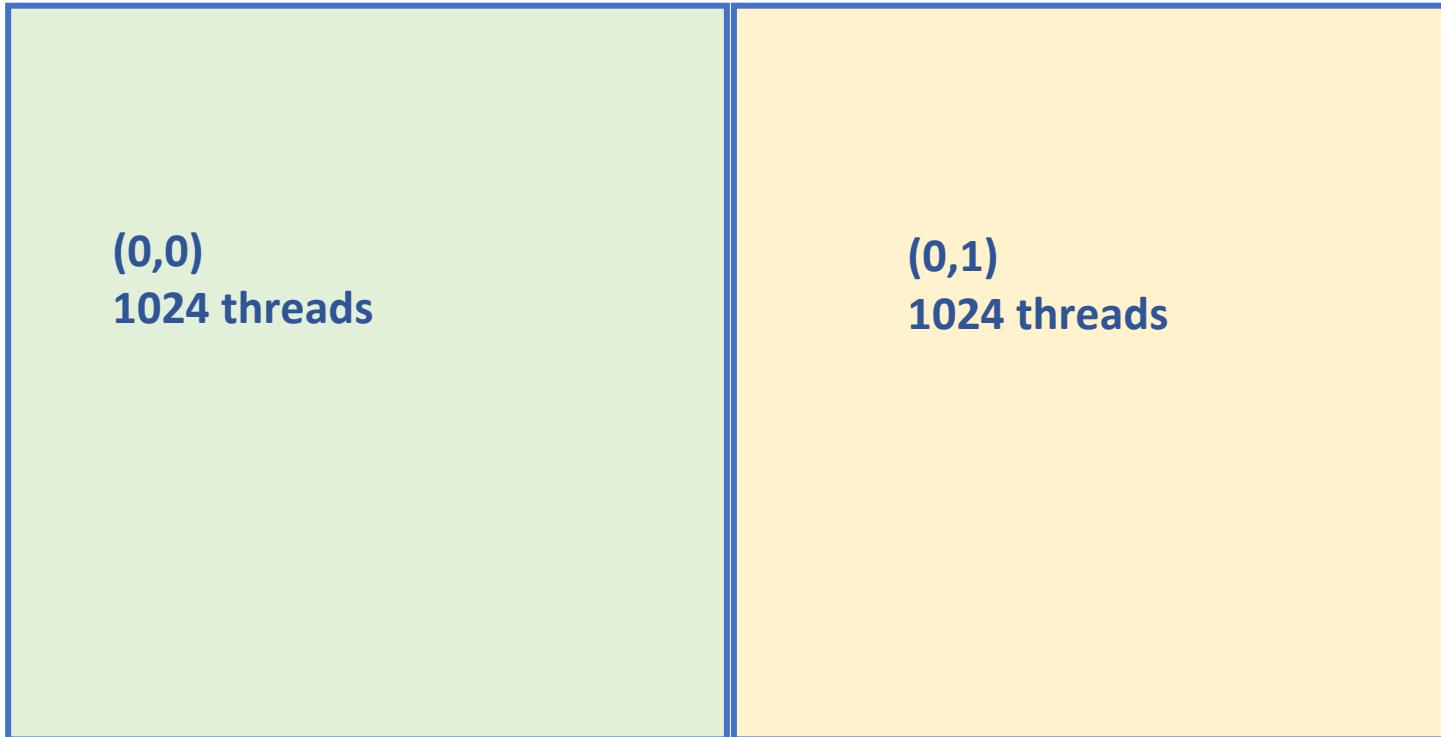
Execution of a block on a GPU

Grid (3x2 blocks)



Execution of a block on a GPU

Grid (2 blocks)



50% usage

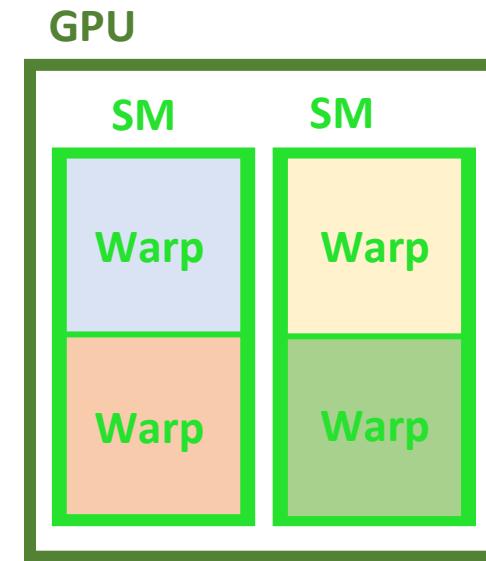
Not enough blocks can result in underutilization!

Execution of a block on a GPU

Multiple blocks can execute in the same SM

This can happen if a block does not have enough warps to fill the SM

But remember: a block is always executed in a single SM.



Outline

Day 1

(09:00 – 13:00)

- ❑ Introduction
- ❑ Coding a CUDA kernel
- ❑ Example
 - ❑ Matrix Multiplication + Exercise
- ❑ Real world implementations
 - ❑ Genomes simulation
 - ❑ Random numbers generation
- ❑ Homework

Outline

Day 1

(09:00 – 13:00)

- ❑ Introduction
- ❑ Coding a CUDA kernel
- ❑ Example
 - ❑ Matrix Multiplication + Exercise
- ❑ Real world implementations
 - ❑ Genomes simulation
 - ❑ Random numbers generation
- ❑ Homework